

# Database Administration

José Orlando Pereira

Departamento de Informática  
Universidade do Minho



# Query processing

a
2
3

"select a from X natural join Y where c = 3;"

X

a	b
1	aaa
2	bbb
3	ccc


Y

b	c
aaa	1
bbb	2
bbb	3
ccc	3
ddd	4


# Compilation

SQL { "select a from X natural join Y where c = 3;"

Relational  
algebra {



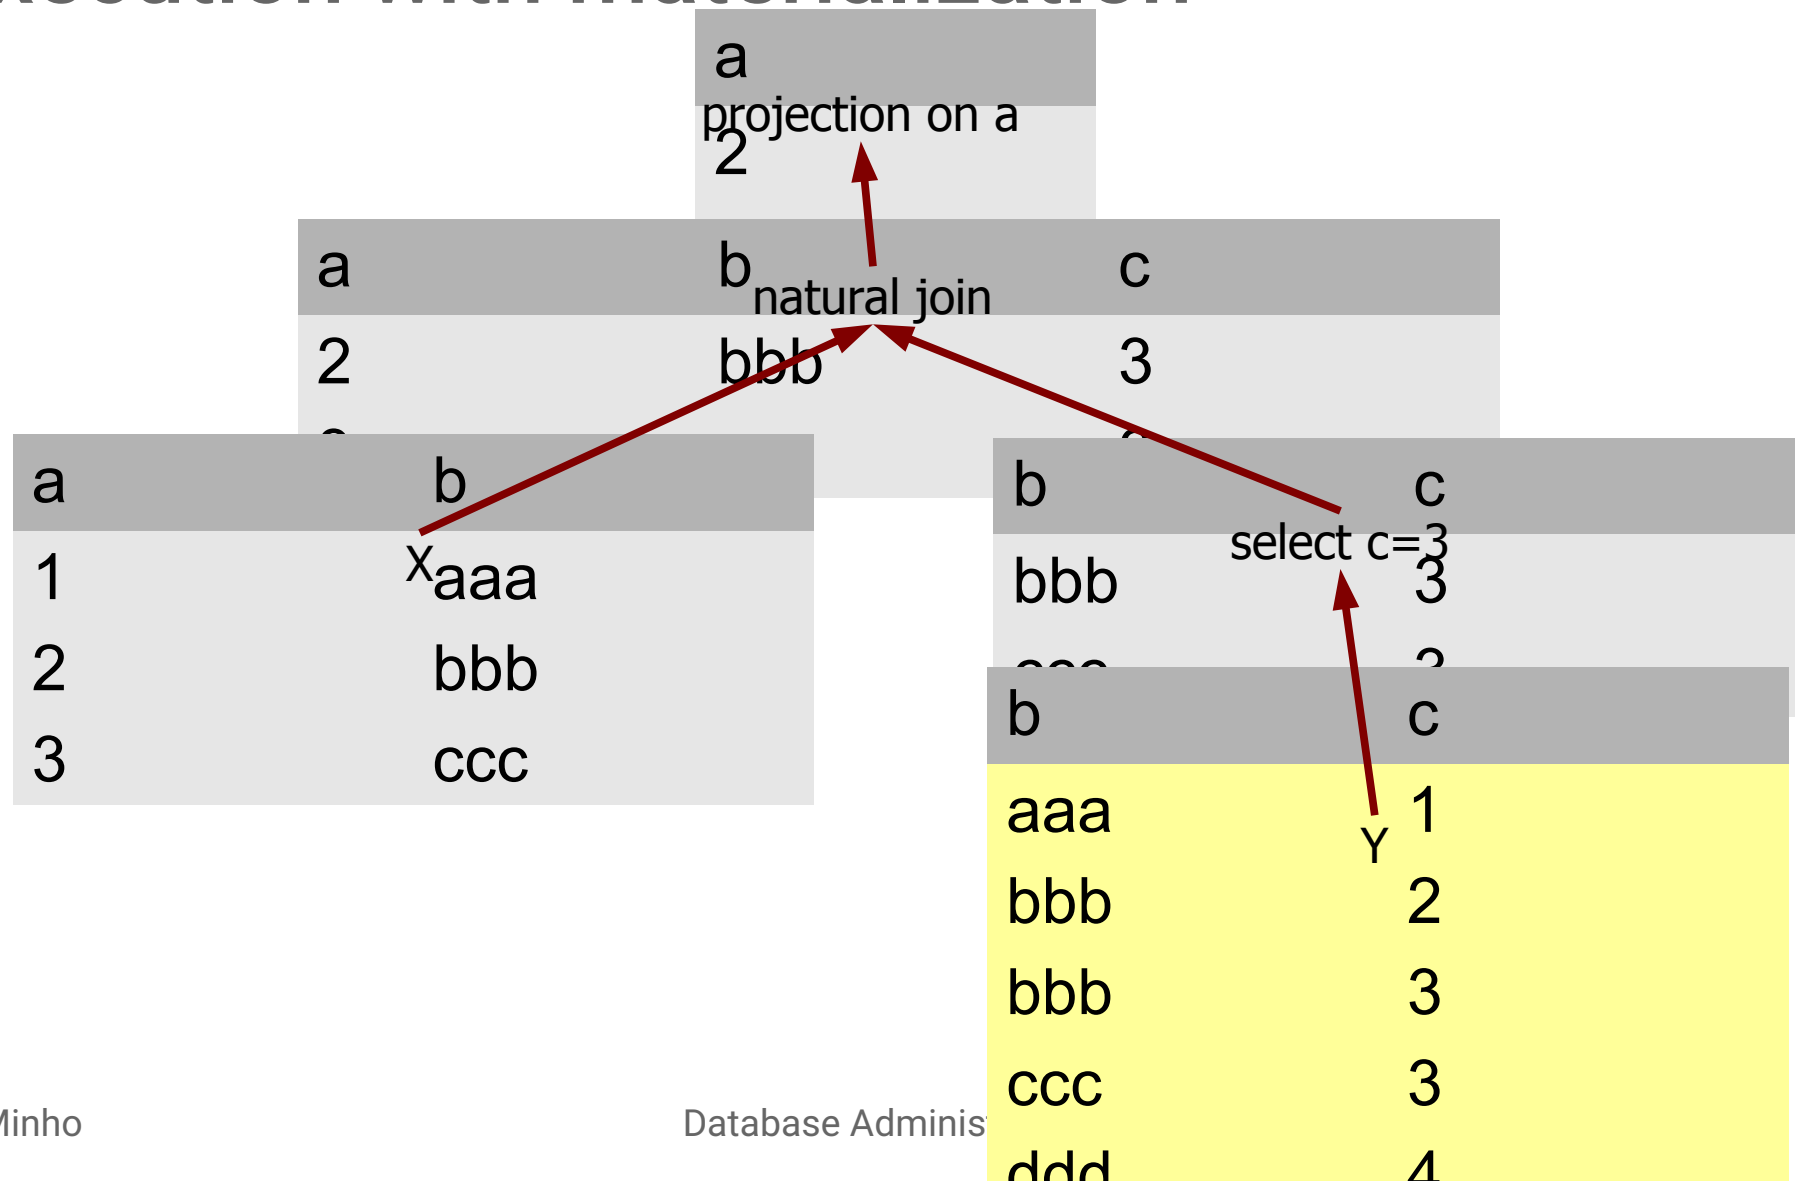
# Logical execution



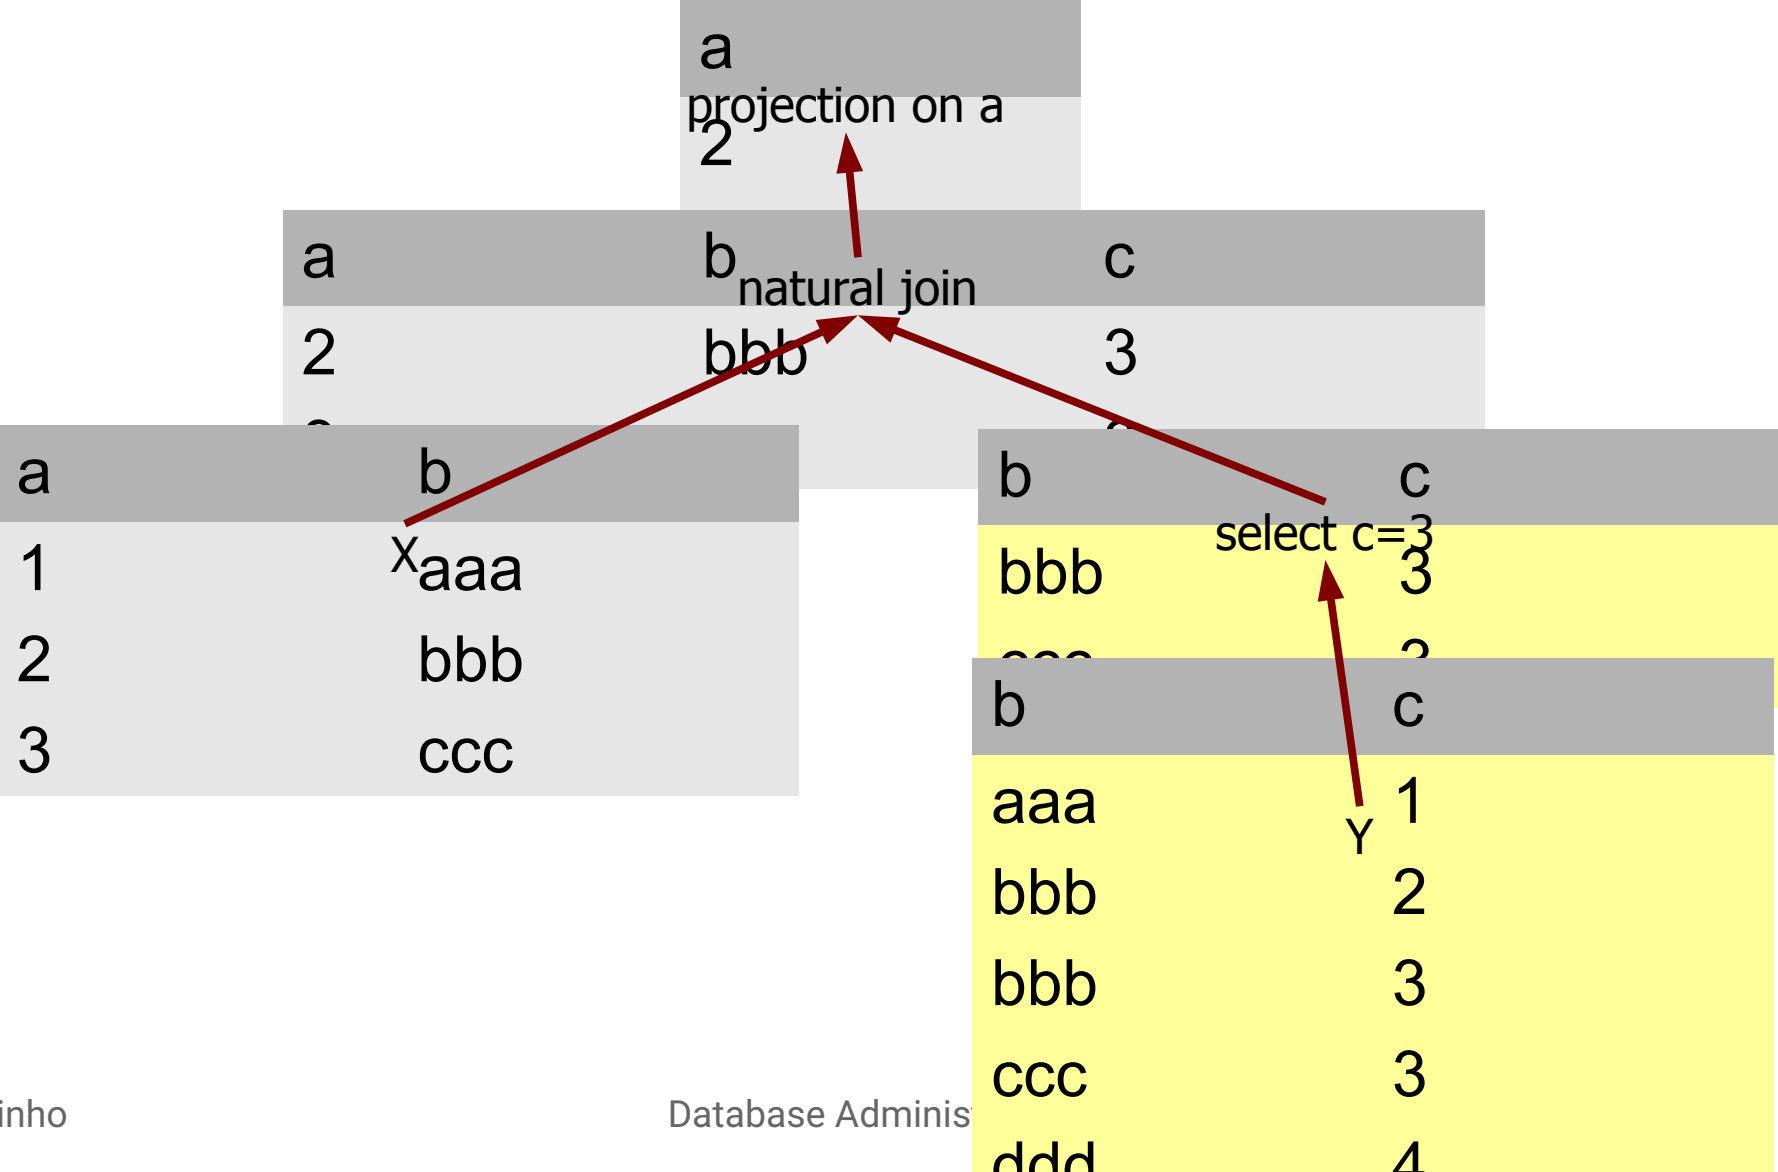
# Materialization

- Each operator is a function:
  - Returns a relation
  - Parameters are other relations (possibly, returned from operators)
- Computation order:
  - From leaves to root

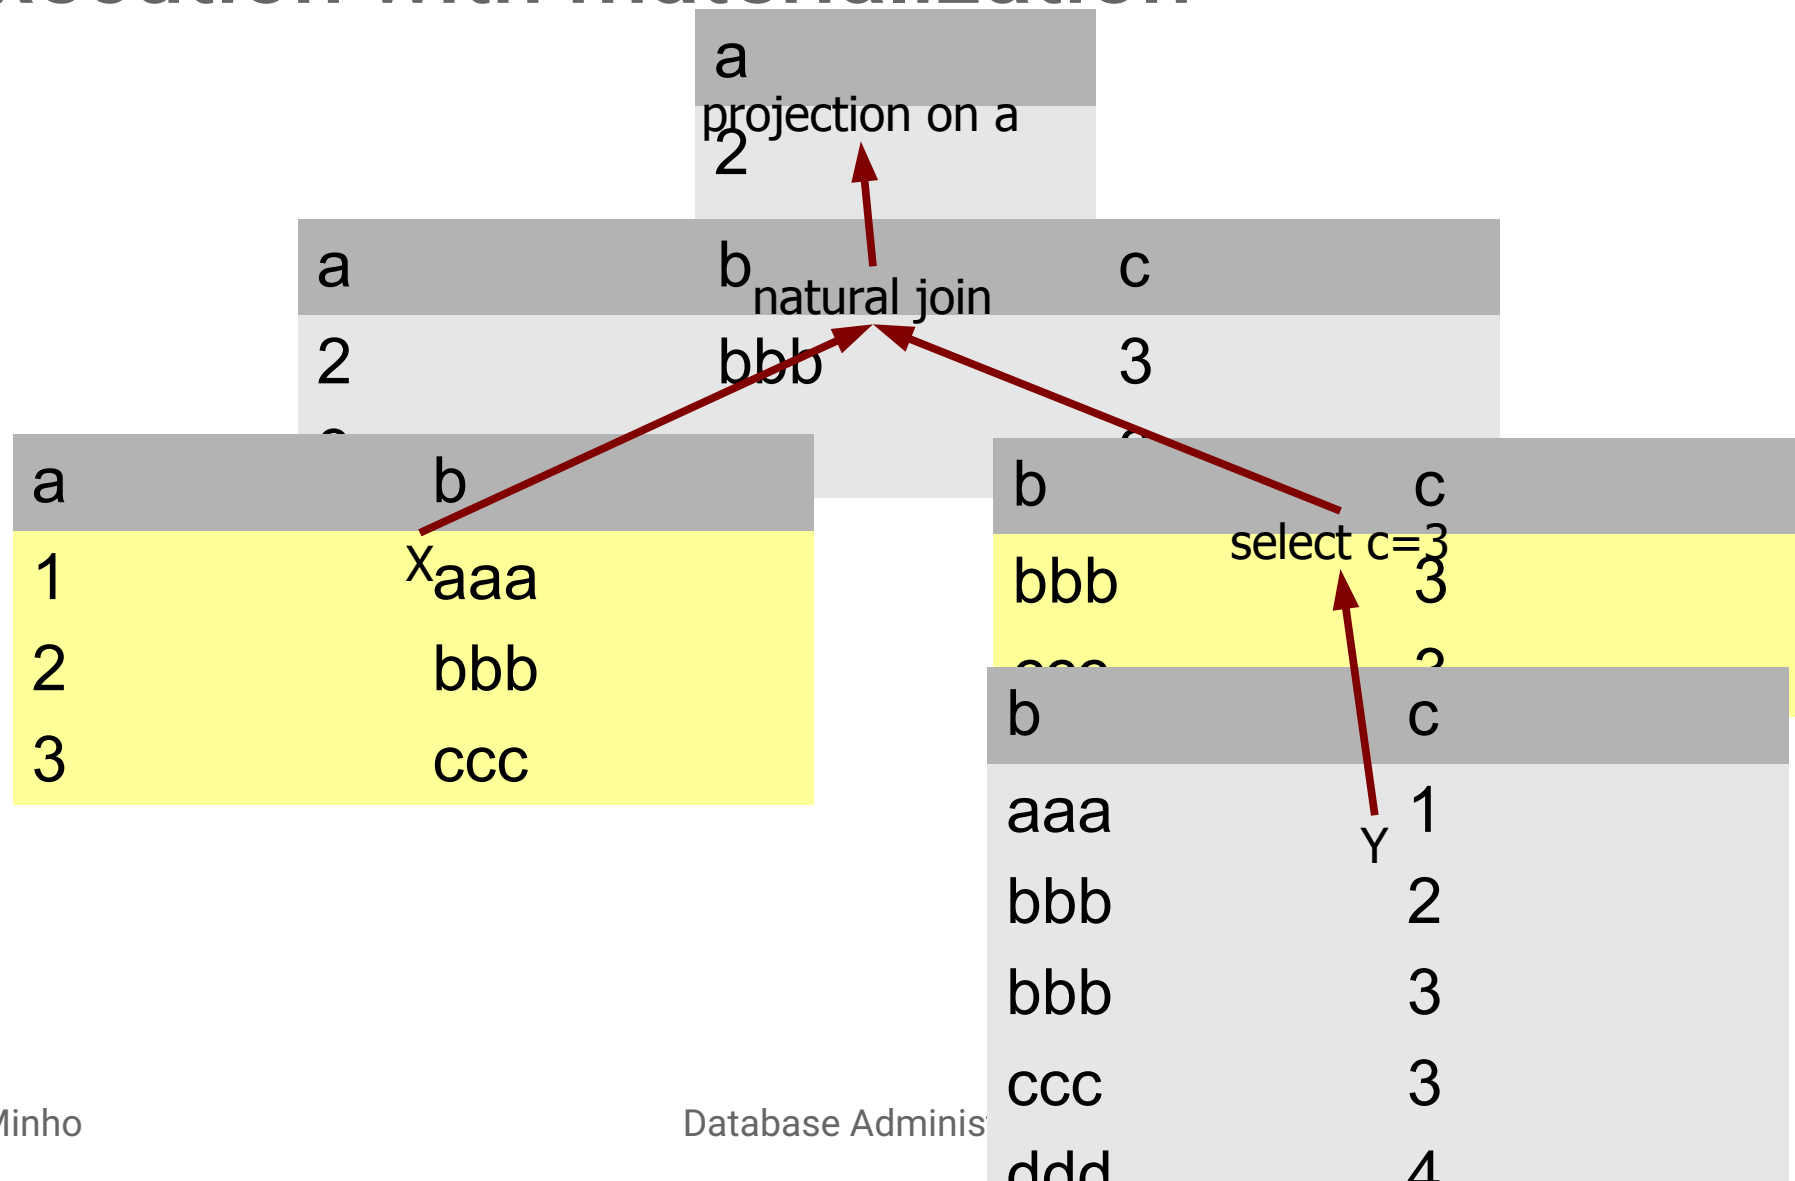
# Execution with materialization



# Execution with materialization

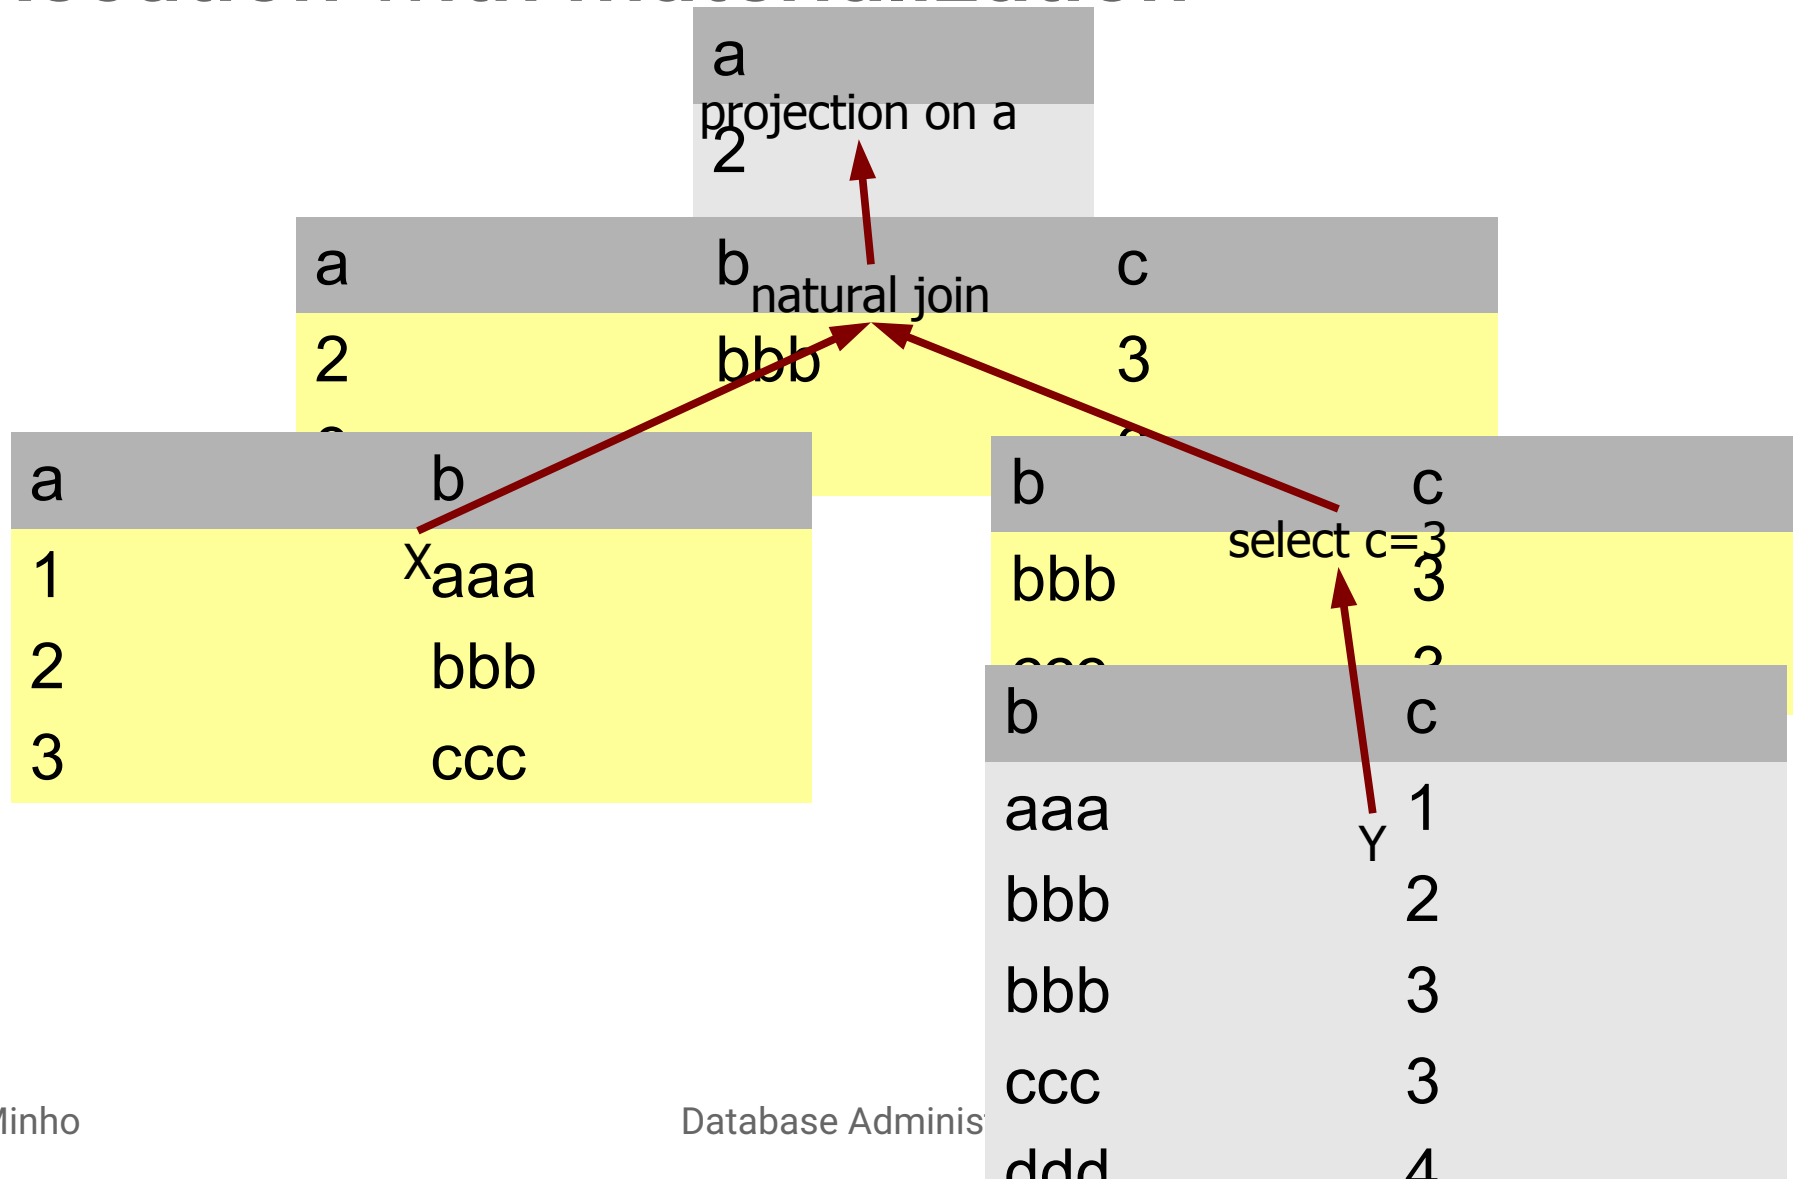


# Execution with materialization






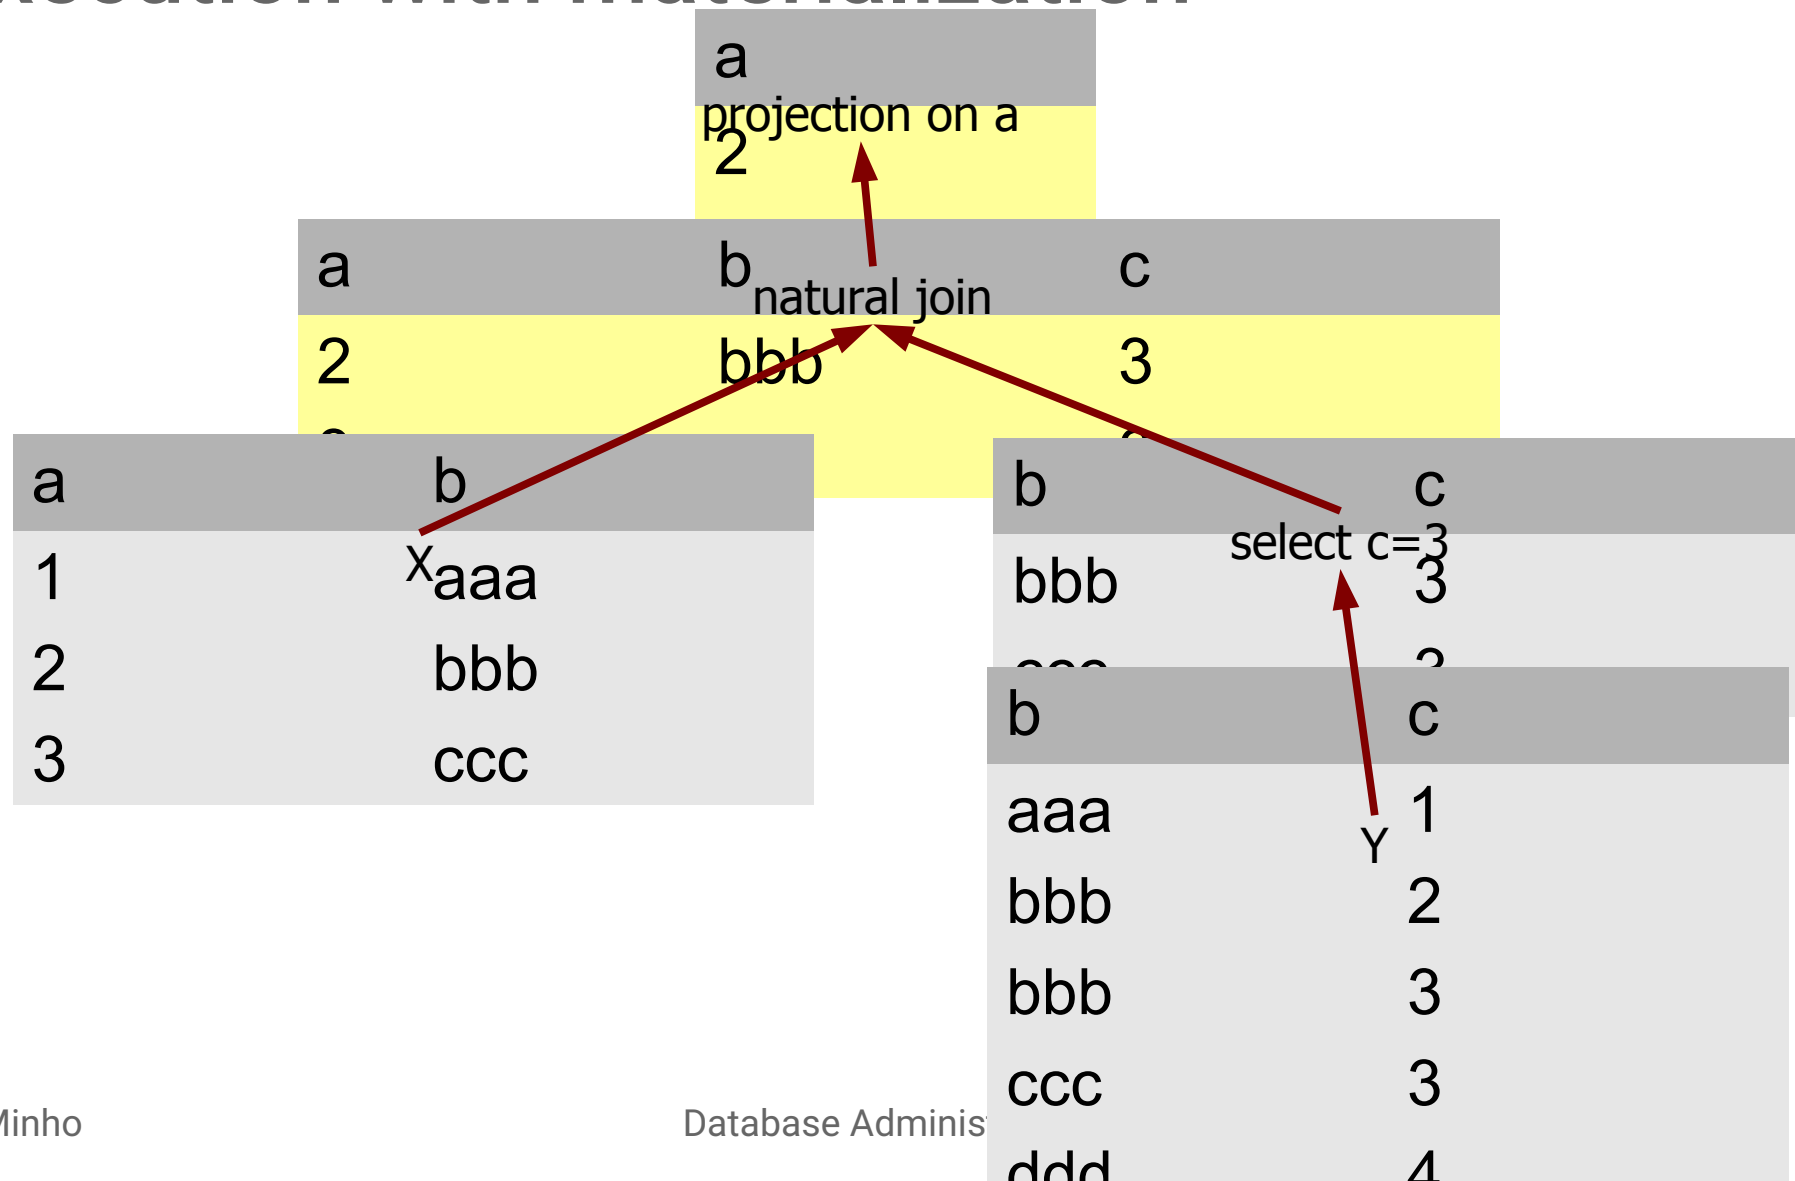
# Execution with materialization




# Execution with materialization



# Execution with materialization



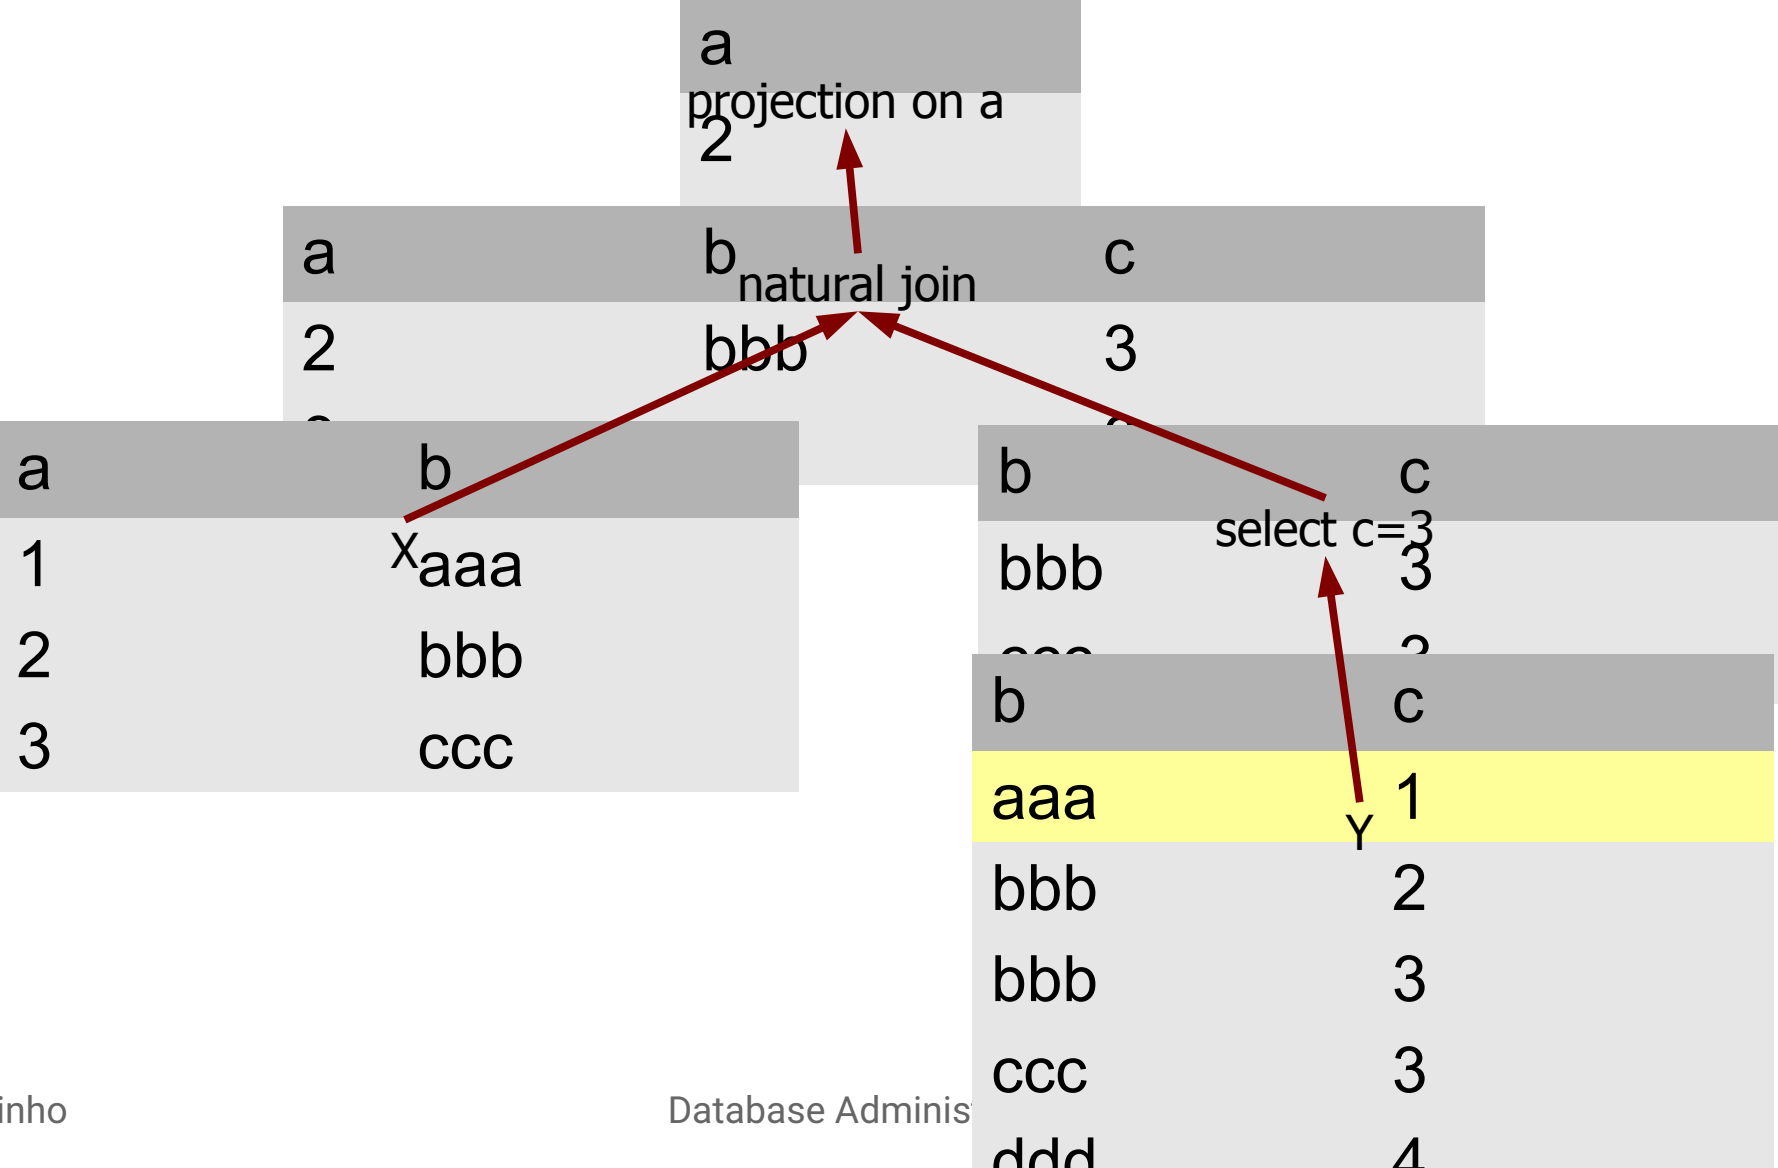
# Execution with materialization



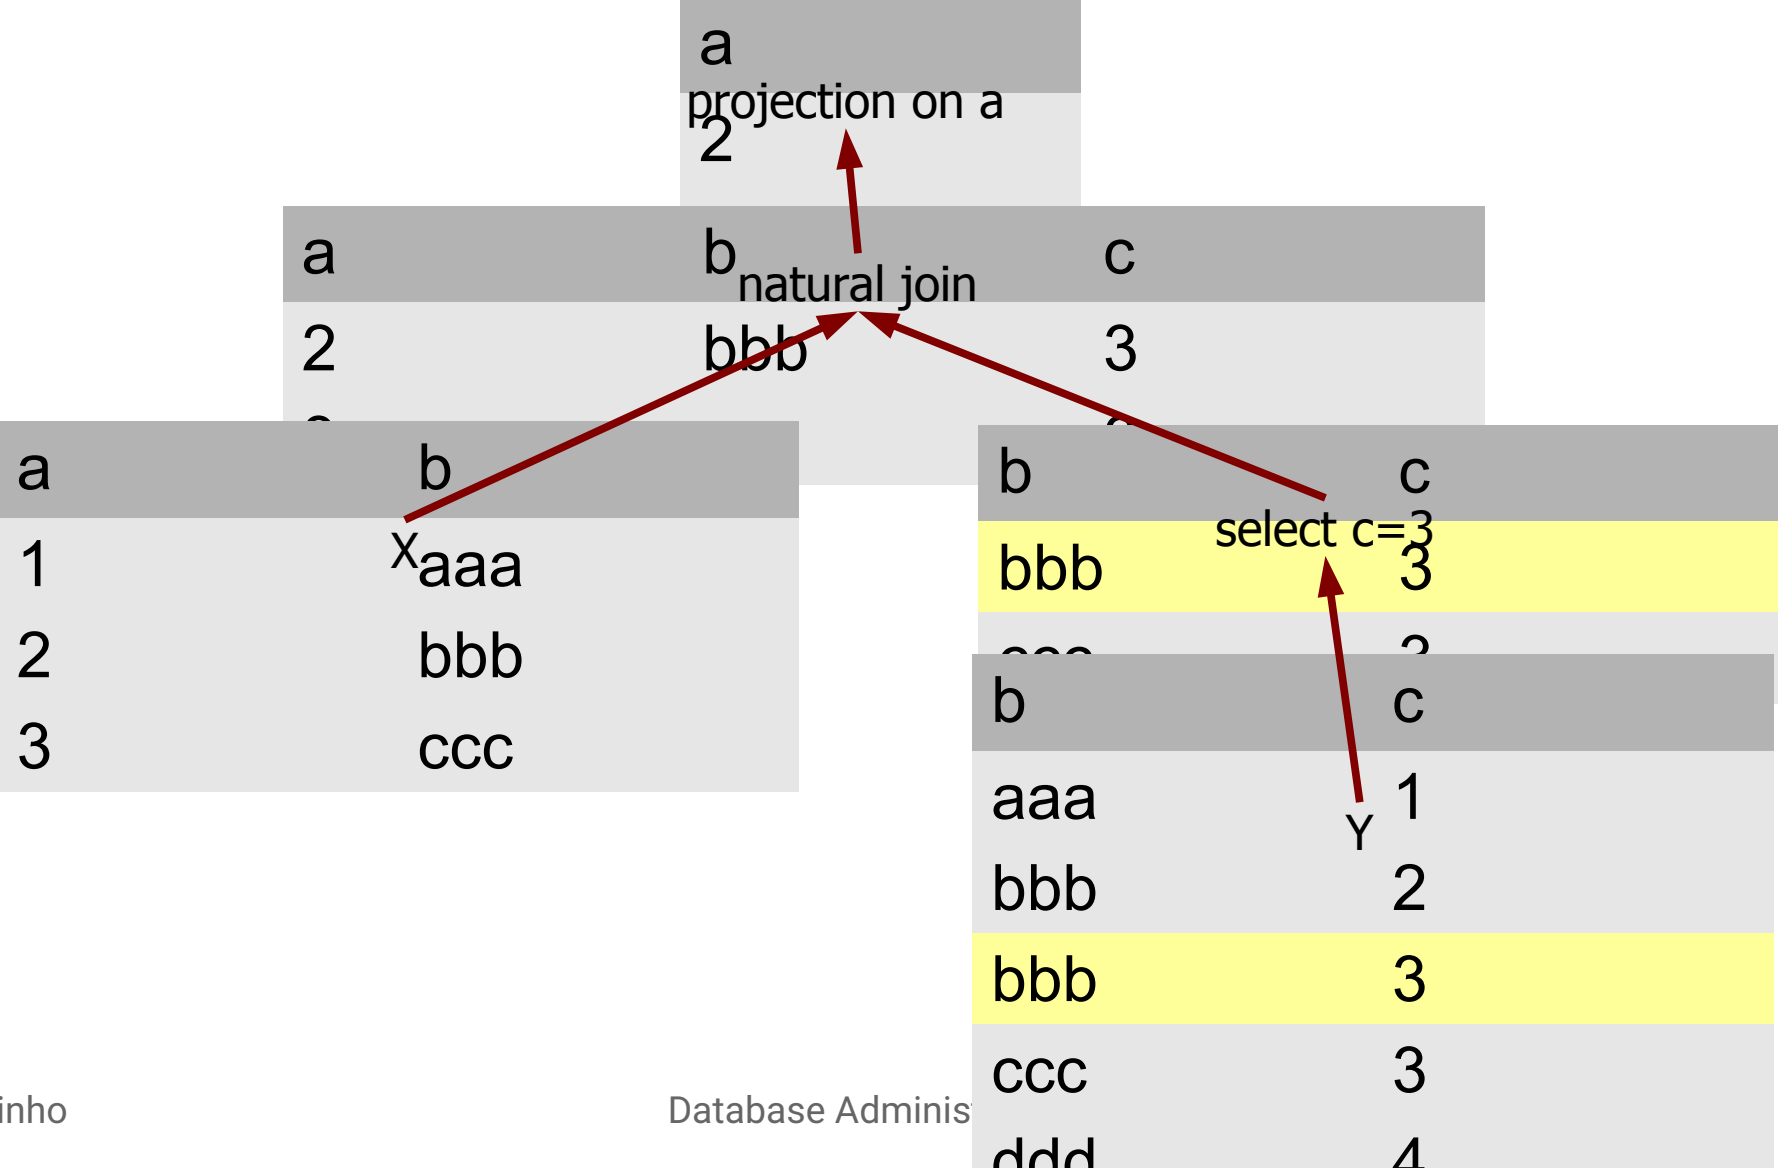
# Iteration

- Each operator is an object:
  - Interface similar to `java.util.Iterator`:
    - `open()` - get ready to return first record
    - `next()` - return next record
    - `close()` - no more records required
  - Constructor parameters:
    - Other operator objects
- Computation order:
  - From leaves to root, for each record


# Execution with iteration



# Execution with iteration

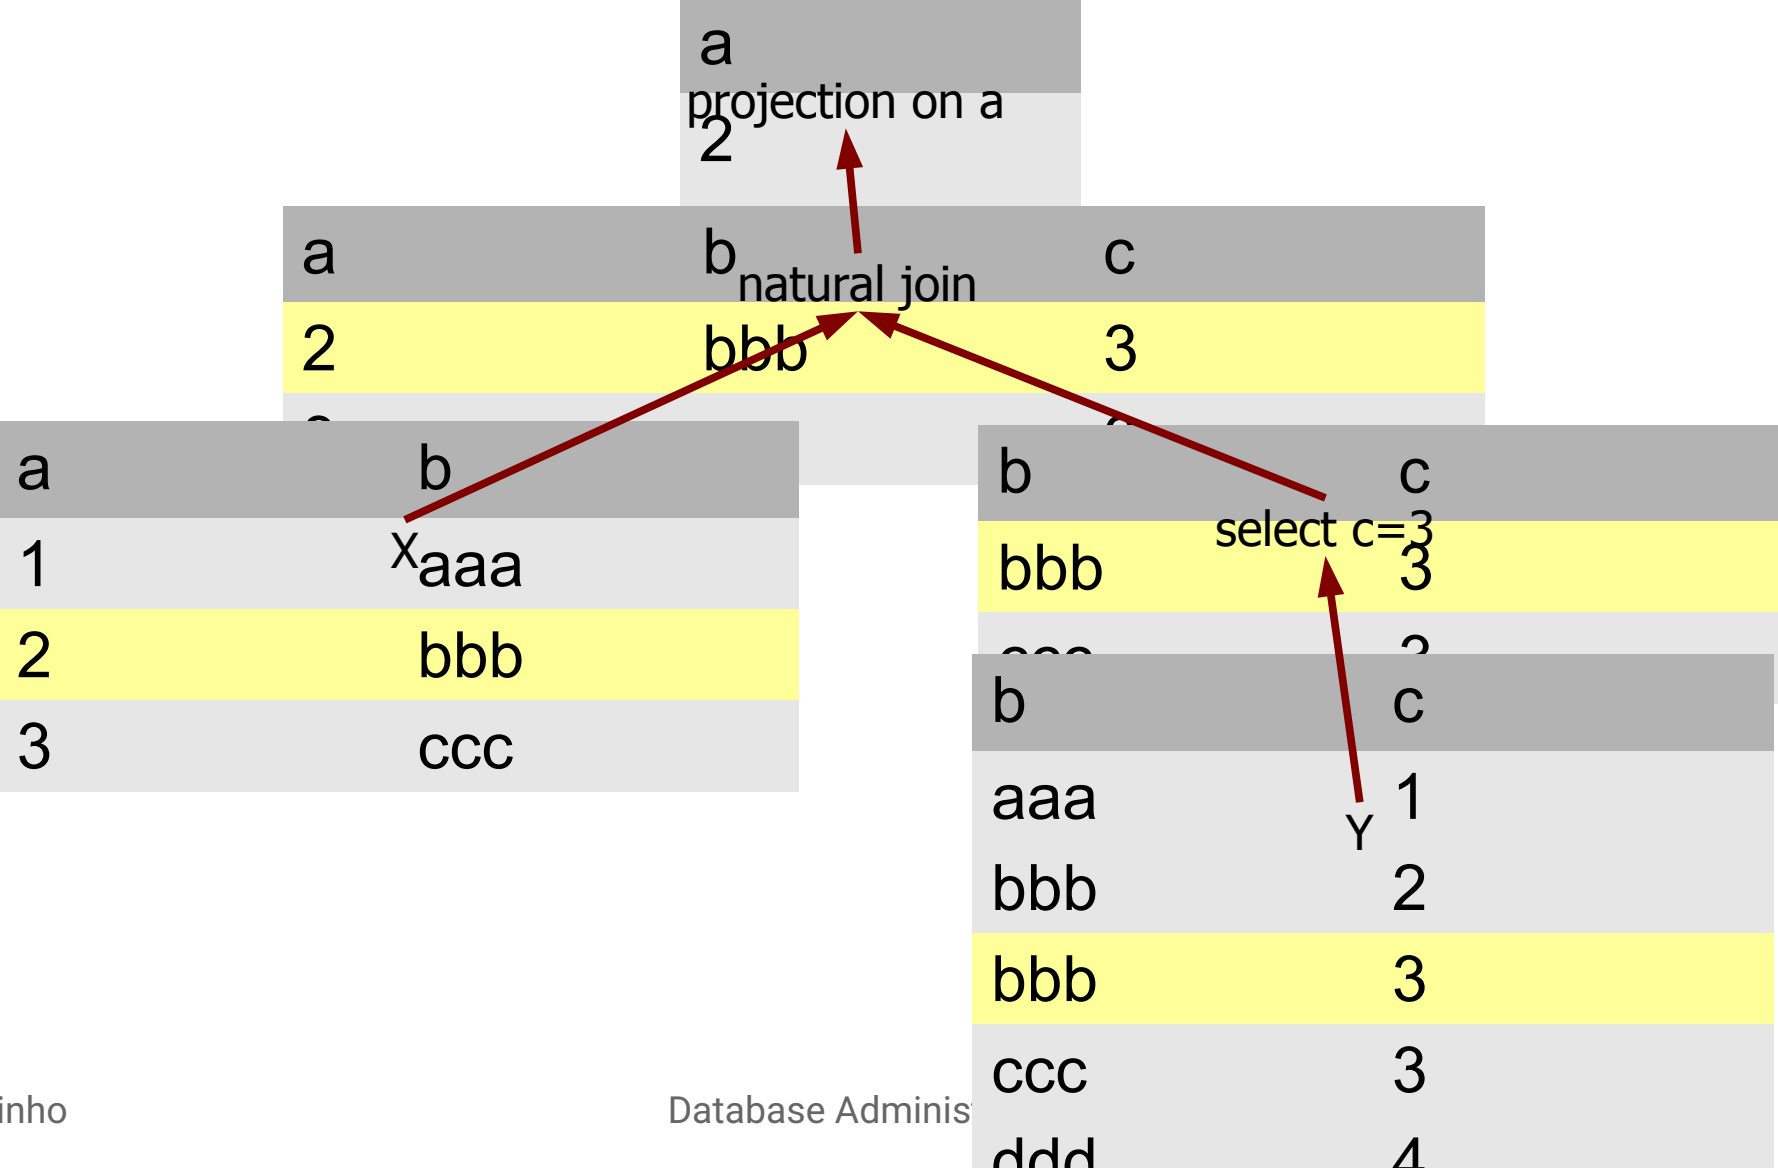


# Execution with iteration

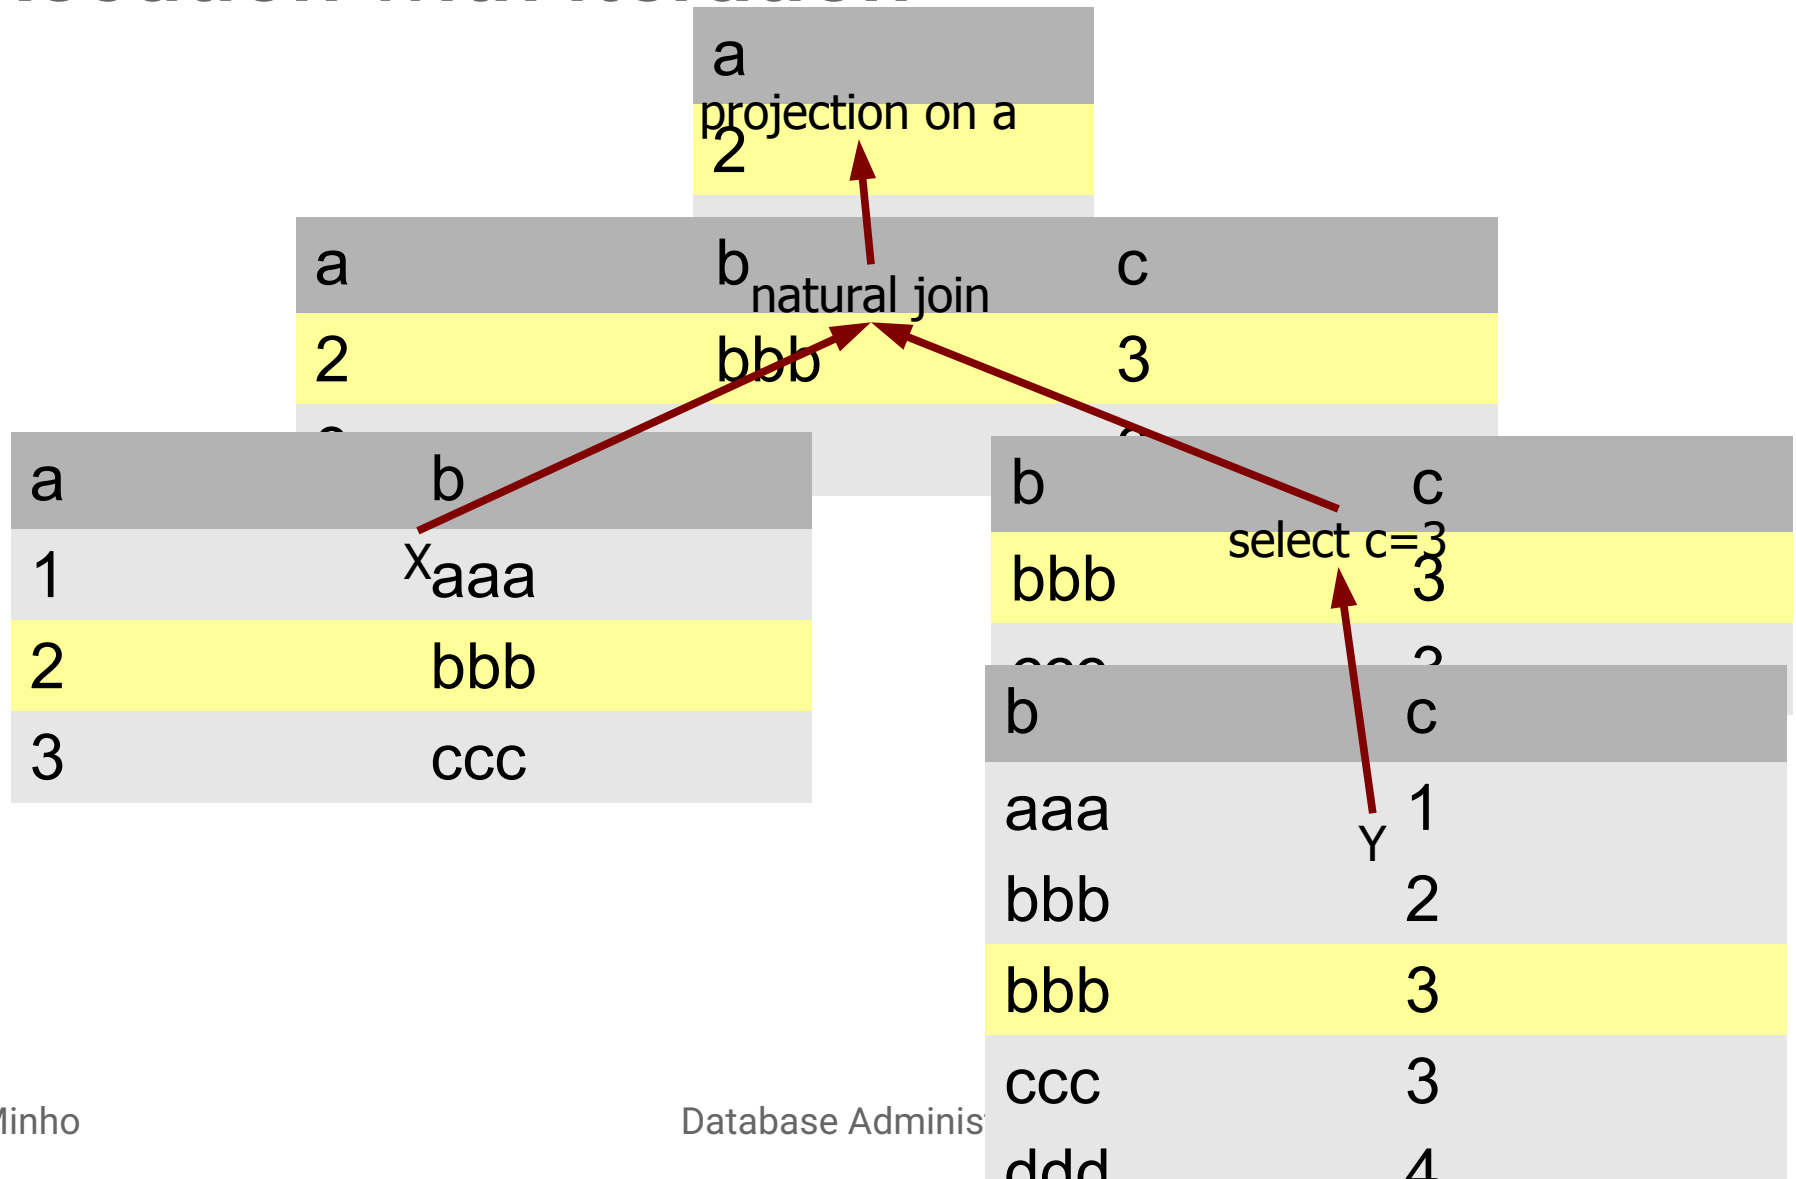





# Execution with iteration



# Execution with iteration



# Execution with iteration



# Materialization vs Iteration

- Iteration avoids caching entire relations
- Materialization avoid reading records more than once
- Can mix both:
  - A materialization operator obtains all records upon first invocation of open
  - Returns records from cached copy on iteration

# Roadmap

- What physical operators exist for each logical operation?
- Later: How are physical operators selected?

# One-pass, record-at-a-time

- Operators:
  - Sequential scan
  - Selection
  - Projection
- Memory requirements:
  - No more than one record required
  - Always possible

# User defined functions (UDFs)

- Functions can be defined in various languages
  - Python
- Scalar functions used in projections/selections:  
`SELECT a, f1(a) FROM t;`  
`SELECT * FROM t WHERE f2(a)`
- Table functions can be used in sequential scans:  
`SELECT * FROM f3(...);`
- User defined functions can access external services:
  - Web services
  - GenAI

# One-pass, full relation, unary

- Duplicate elimination:
  - Cache unique records
  - “select distinct \* from X;”
- Grouping and aggregation:
  - Cache groups
  - “select count(\*) from X group by b;”
- Sorting:
  - Cache all records and sort in memory
  - “select \* from X order by b;”




# One-pass, full relation, binary

- Union, difference, intersection, product, join:
  - Read and cache the smallest relation
  - Organize for fast look-up (e.g. hash)
  - Read and operate on each record from the largest relation

# One-pass, full relation, binary

- Load smaller table into memory and add search structure:

a	b	c
2	bbb	3
3	ccc	3

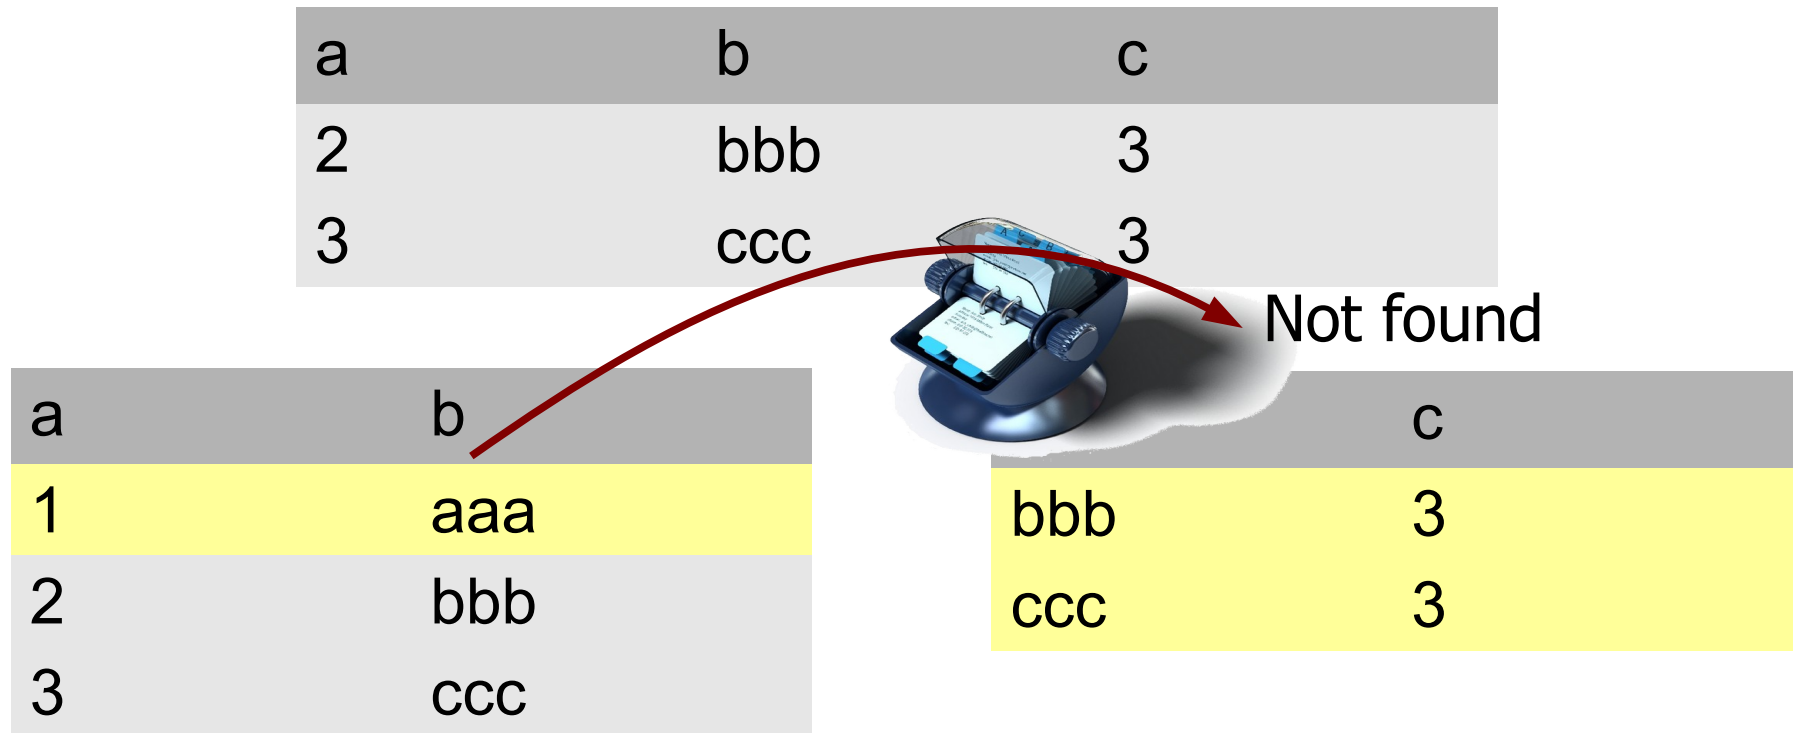


a	b
1	aaa
2	bbb
3	ccc

c
bbb 3
ccc 3


# One-pass, full relation, binary

- Test each record from the largest relation:



# One-pass, full relation, binary

- Test each record from the largest relation:



# Nested-loop join (NLJ)

a	b
1	aaa
2	bbb
3	ccc

b	c
bbb	3
ccc	3
b	c

a	b
1	aaa
2	bbb
3	ccc

b	c
bbb	3
ccc	3
b	c
bbb	3
ccc	3

# Nested-loop join (NLJ)

- Memory requirements:
  - One record from each relation
- Operations:
  - If outer loop has N records
  - Reads inner relation N times

# Block-based NLJ

- Much smarter: Execute NLJ by blocks
- Memory requirements:
  - One block from each relation
- Operations:
  - If outer loop has  $N$  records /  $B$  blocks ( $B \ll N$ )
  - Reads inner relation  $B$  times ( $B \ll N$ !)

# Case study

- Tables:
  - Client: Id, Name, Address, Data<sup>(\*)</sup>
  - Product: Id, Description, Data<sup>(\*)</sup>
  - Invoice: Id, ProductId, ClientId, Data<sup>(\*)</sup>
- Pre-populate Client and Product with  $2^n$  items

<sup>(\*)</sup> Strings with arbitrary data...



# Case study

- Sell:
  - Add invoice record
- Account of a specific client:
  - names of items sold to that client
- Top 10 products:
  - 10 most sold products
- Generate client and product ids with:  
`rand.nextInt(MAX) | rand.nextInt(MAX)`

# Benchmarking

- Repeat workload for a variable number of client threads
- Discard initial and final periods
- Measure:
  - Response time (duration of transactions)
  - Throughput (rate of execution)

