

Database Administration

José Orlando Pereira

Departamento de Informática
Universidade do Minho



Motivation

- Problem:
 - select x from Y
where z = 'k';
- Plan:

select z = 'k'

scan Y

- Cost?


z	x
d	2
c	6
g	2
k	2
h	3
a	1
b	1
f	3
d	2
k	1
j	2
l	1
...	...

Index

- Makes it easy to find pages containing interesting data
- Smaller than data
 - Fits in memory?
- Efficient look-up:
 - Identity (=)
 - Ranges
 - LIKE
 - ...

z	x
d	2
c	6
g	2
k	2
h	3
a	1
b	1
f	3
d	2
k	1
j	2
l	1
...	...

B-Tree



B-Tree

- Insert:
 - If free entry not available, split leaf
 - Recursively insert new leaf in upper layer
 - Tree grows towards the root!
 - Add entry to leaf
- Delete:
 - Remove entry from leaf
 - If enough space available, collapse leafs
 - Recursively delete leaf in upper layer


B-Tree

- Desirable characteristics:
 - Balanced
 - $\log(n)$ depth
 - Fit for block I/O
- Supports:
 - Identity look up
 - Range queries / Ordered scan
 - Updates

Composite indexes

- Index on (X,Y):
 - Answers equality on (X,Y)
 - Answers equality and interval on X alone
 - Answers equality on X and interval on Y
- Index on expression, e.g. X+Y
 - Answers equality and interval on X+Y

Dispersion




Clustered indexes


- Problem:
 - $\#blocks \gg (\#records / \text{records per block})$
 - Read each block multiple times
- A clustered index:
 - Records are (roughly) sorted according to the index
 - No sorting within a block is needed
 - Free space may be kept for insertions

Multi-criteria filtering

- Problem:
 - select x from Y where z = 'k' and x = 4;



Z	X
d	2
c	6
g	4
h	2
h	3
k	4
b	1
f	3
d	2
k	1
j	2
l	1
...	...






Multi-criteria filtering

- Sequential scan
- Either X or Z plus scan
- Composite index on (Z,X):
 - Often, many columns and combinations
- Typical examples:
 - Search with a combination of features
 - Hotel booking, real estate, on-line shopping, ...
- What if “or” instead of “and”?
 - (T2 or T3) and in Braga


Bitmap indexes

- Step 1:
 - Build bitmap for each clause



Bitmap indexes

- Step 2:
 - Combine with logical operators



Bitmap indexes

- Step 3:
 - Traverse bitmap and read table


... and ...	Z	X
0	d	2
0	c	6
0	g	4
0	h	2
0	h	3
1	k	4
0	b	1
0	f	3
0	d	2
0	k	1
0	j	2
0	l	1
...

Motivation

- Problem:

```
select z from Y  
order by v <->  
      (select v from Y where z='k')  
limit 10;
```

- Plan:




- Cost?

z	v
d	[.9, ...]
c	[.6, ...]
g	[.7, ...]
k	[.2, ...]
h	[.3, ...]
a	[.4, ...]
b	[.5, ...]
f	[.3, ...]
d	[.5, ...]
x	[.1, ...]
j	[.9, ...]
i	[.8, ...]
...	...


Vector indexes

	v
z	
d	[.9, ...]
c	[.6, ...]
g	[.7, ...]
k	[.2, ...]
h	[.3, ...]
a	[.4, ...]
b	[.5, ...]
f	[.3, ...]
d	[.5, ...]
x	[.1, ...]
j	[.9, ...]
l	[.8, ...]
...	...

- Graph structure by proximity in multi-dimensional space:
 - Points to original rows



Vector indexes



- Problem:
`select z from Y
order by v <->
(select v from Y where z='k')
limit 10;`
- Breadth-first traversal of graph from 'k' produces rows in increasing distance
- Simple plan:

limit
↑
index scan

Motivation


- Assumptions:
 - Several TB of data
 - ~50%, $y=1$
 - ~50%, $y=2$
 - a few, $y=3$

z	y
d	1
c	2
g	1
k	2
h	3
a	1
b	1
f	2
d	2
k	1
j	2
l	1
...	...

Motivation

- Problem:
 - `select count(*) from X where y = 1;`

- Possible plans:



- Cost?

z	y
d	1
c	2
g	1
k	2
h	3
a	1
b	1
f	2
d	2
k	1
j	2
l	1
...	...

Motivation


- Keep results cached
- Update when needed

select y, count(*) from X group by y


y	count
1	773647263
2	765732332
3	1

z	y
d	1
c	2
g	1
k	2
h	3
a	1
b	1
f	2
d	2
k	1
j	2
l	1
...	...

Views



Materialized views



Maintaining materialized views

- Periodically run the query and update the view
- Update the view when data changes

DIY Materialized Views

- Updating with AFTER triggers:

```
-- Create view
select sum(value) into mv_sum_items from items;

-- Update view
create function upd_sum_items() returns trigger as '
BEGIN
    update mv_sum_items set sum = sum + new.value - old.value;
    return new;
END
language 'plpgsql';

create trigger upd_sum
after update on items
for each row execute procedure upd_sum_items();
```

Using materialized views

- Automatically used by the planner:
 - Indexed views in MS SQL Server
- Used explicitly in queries:
 - Materialized views in Oracle
 - DIY materialized views everywhere
 - Developer tip:
 - Using views allows the DBA to select which ones to materialize

Conclusions

- Indexes and mat. views = Redundancy!
- Trade-off between:
 - Complexity of operationsand:
 - Disk space used
 - Usage of main memory
 - Effort when updating
- Usefulness depends on workload mix