

# Database Administration

José Orlando Pereira

Departamento de Informática  
Universidade do Minho




# Roadmap

- What physical operators exist for each logical operation?
- How are physical operators implemented and composed?
- Later: How are physical operators selected?


# Execution with materialization

- Each operator is a function:
  - Returns a relation
  - Parameters are other relations (possibly, returned from operators)
- Computation order:
  - From leaves to root


# Execution with materialization




# Execution with materialization




# Execution with materialization




# Execution with materialization




# Execution with materialization



# Execution with materialization



# Execution with materialization



# Consequences

- Large amount of intermediate results that need to be stored
  - Might not fit completely in memory
- Potentially wasted work
  - e.g. SELECT ... LIMIT 10


# Challenges

- How to minimize data movement while executing?
- How to achieve the best performance given current computer architectures?


# Iteration (pull model)

- Each operator is an object:
  - Interface similar to `java.util.Iterator`:
    - `open()` - get ready to return first record
    - `next()` - return next record
    - `close()` - no more records required
  - Constructor parameters:
    - Other operator objects
- Computation order:
  - From leaves to root, for each record


# Execution with iteration




# Execution with iteration




# Execution with iteration




# Execution with iteration



# Execution with iteration



# Execution with iteration



# Consequences

- Not applicable to operators that must observe all rows before knowing what is the first to output
  - ORDER BY
  - GROUP BY on an unsorted input
  - ...
  - In these cases, the operator executes all the underlying plan on open() or the first invocation to next()
- Minimizes work with LIMIT clause