

Database Administration

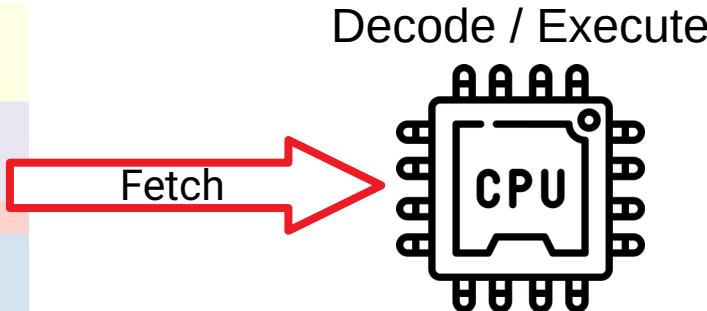
José Orlando Pereira

Departamento de Informática
Universidade do Minho



A model of computing

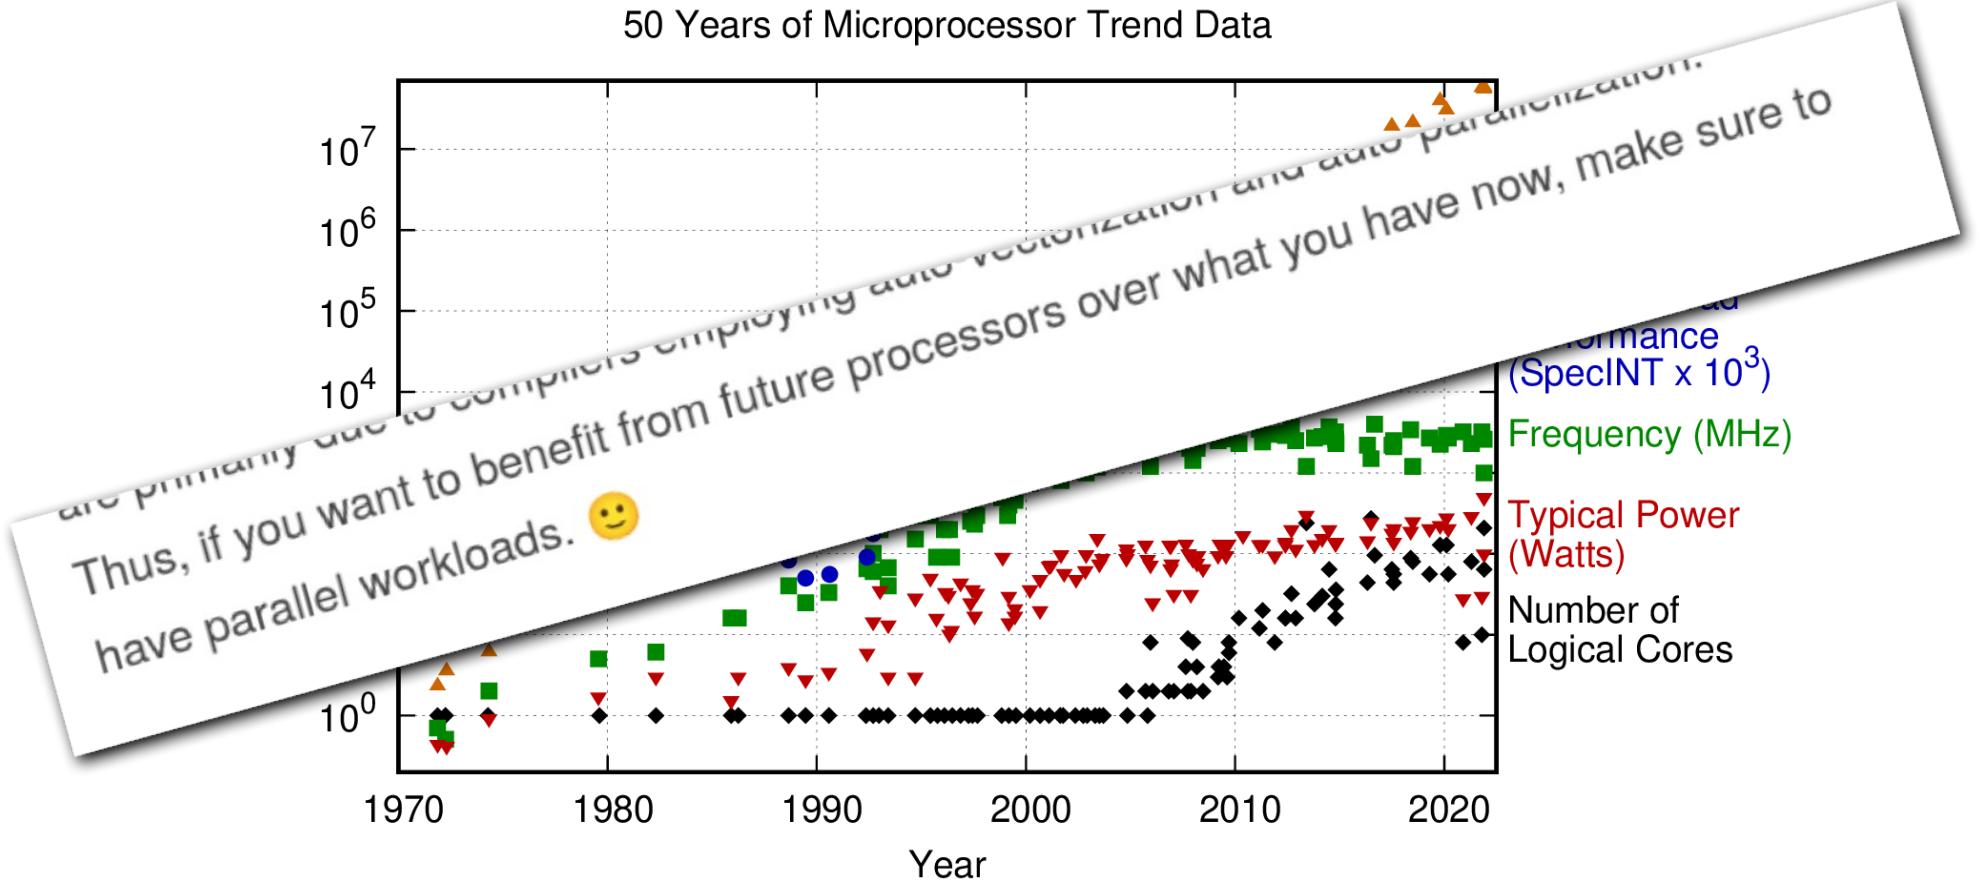
```
5      mov    eax, DWORD PTR [rbp-20]
6      imul   eax, eax
7      mov    DWORD PTR [rbp-8], eax
8      mov    eax, DWORD PTR [rbp-20]
9      add    eax, eax
10     IP -> mov    DWORD PTR [rbp-12], eax
11     mov    DWORD PTR [rbp-4], 5
12     mov    eax, DWORD PTR [rbp-4]
13     sub    eax, DWORD PTR [rbp-12]
14     test   eax, eax
15     jns   .L2
16     mov    eax, DWORD PTR [rbp-12]
17     add    DWORD PTR [rbp-4], eax
18     .L2:
19     mov    edx, DWORD PTR [rbp-4]
20     mov    eax, DWORD PTR [rbp-8]
21     add    edx, eax
22     mov    eax, DWORD PTR [rbp-12]
```



- Challenge for data-intensive programs:
 - Computation is not fast enough

Moore's Law

50 Years of Microprocessor Trend Data

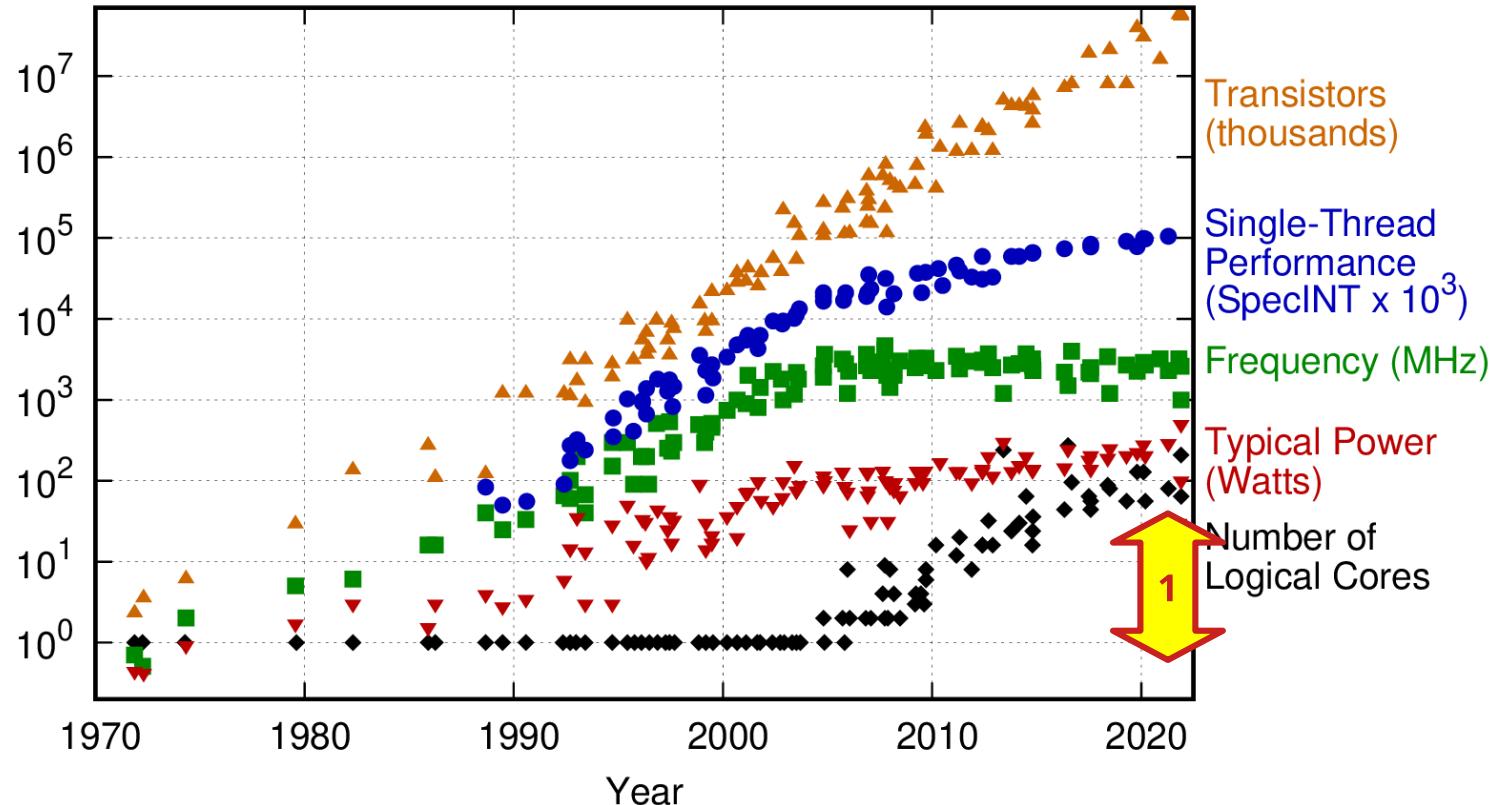


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp

Source <https://github.com/karlrupp/microprocessor-trend-data>

Multiple cores

50 Years of Microprocessor Trend Data

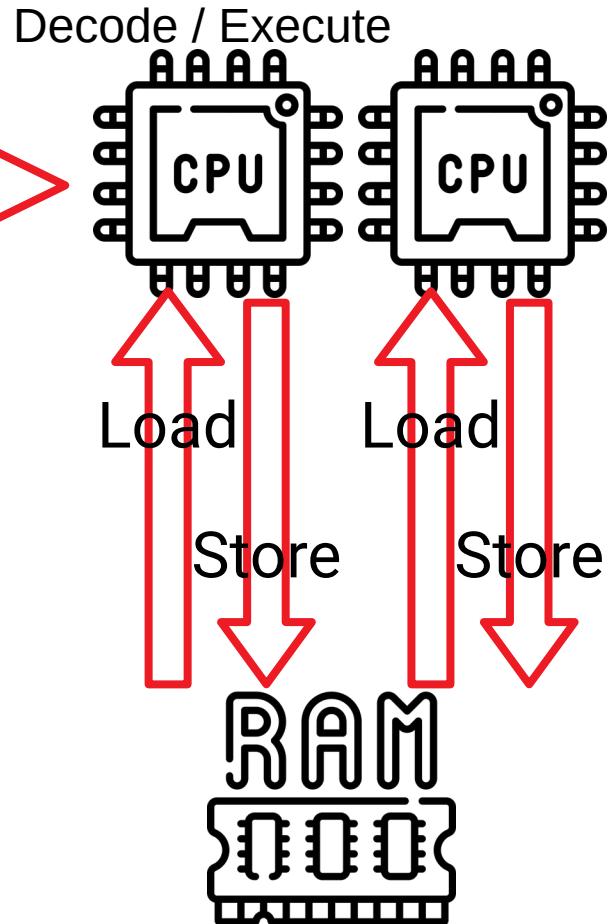


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp

Source <https://github.com/karlrupp/microprocessor-trend-data>

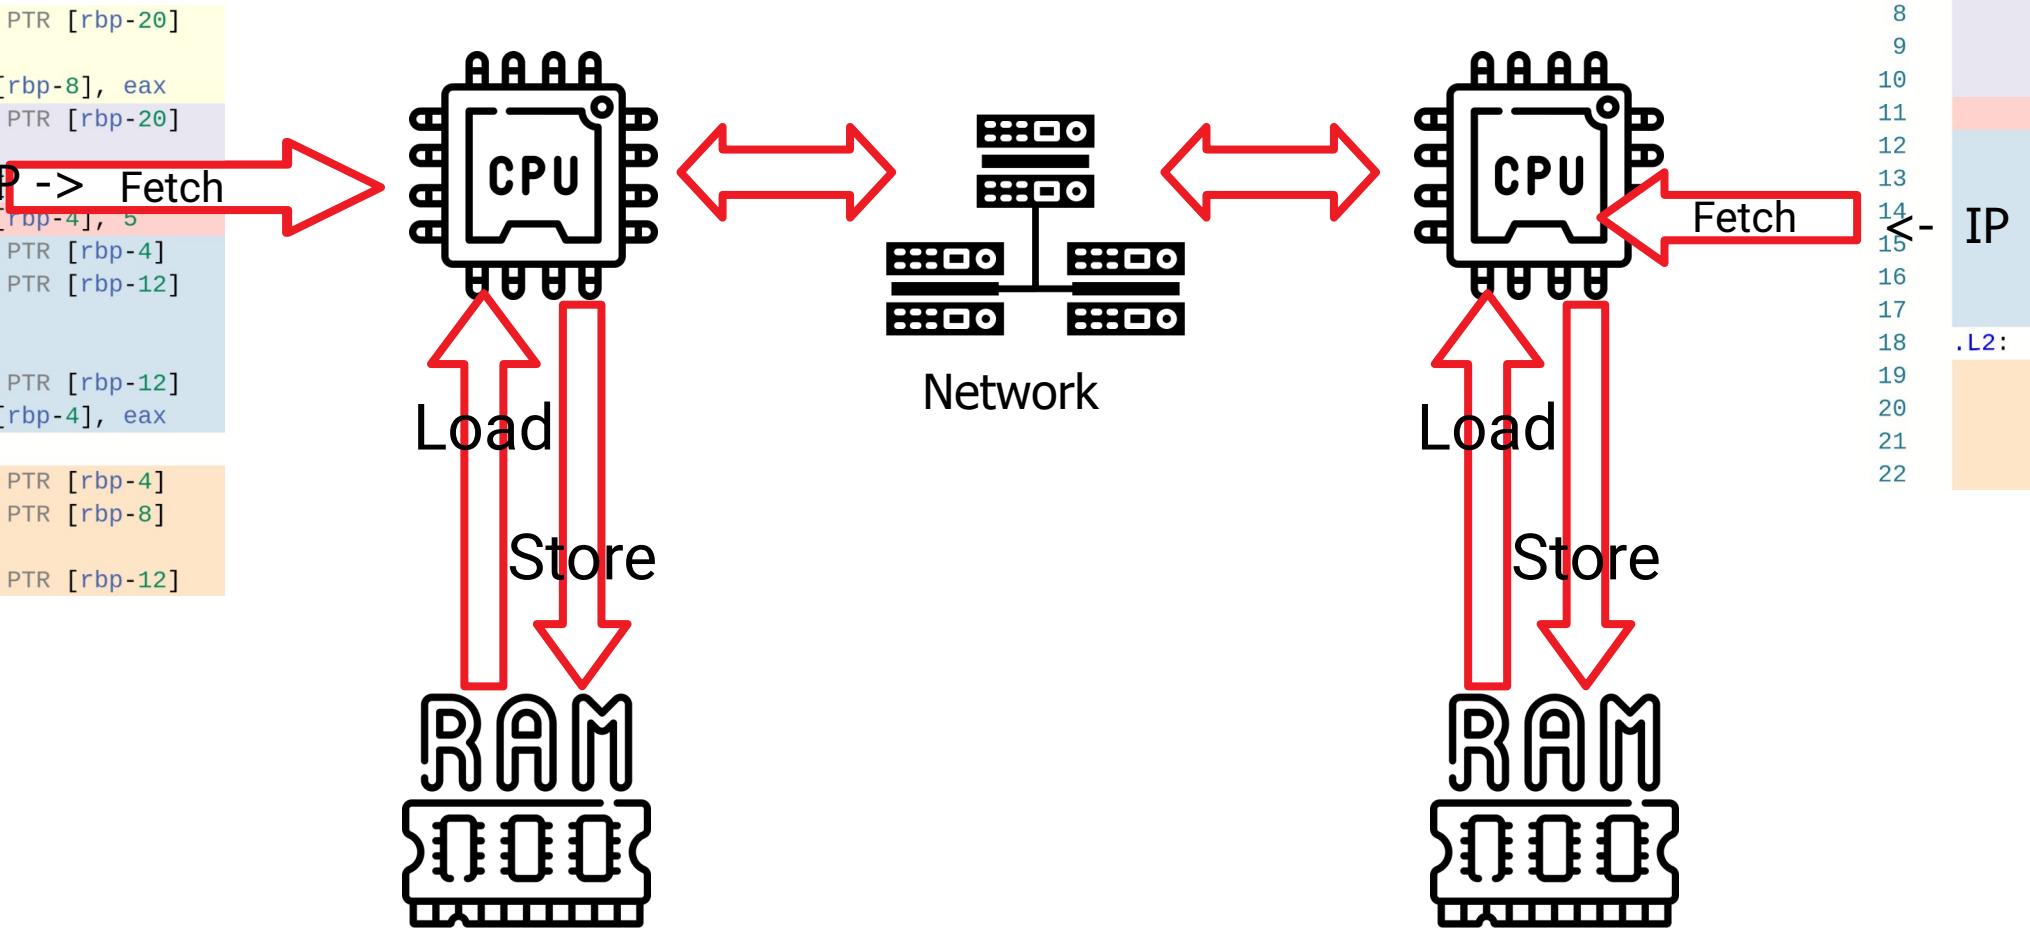
Multi-core

```
eax, DWORD PTR [rbp-20]
eax, eax
DWORD PTR [rbp-8], eax
eax, DWORD PTR [rbp-20]
eax, eax
DWORD PTR IP->[rbp-12], eax
DWORD PTR [rbp-4], 5
eax, DWORD PTR [rbp-4]
eax, DWORD PTR [rbp-12]
eax, eax
L2
eax, DWORD PTR [rbp-12]
DWORD PTR [rbp-4], eax
edx, DWORD PTR [rbp-4]
eax, DWORD PTR [rbp-8]
edx, eax
eax, DWORD PTR [rbp-12]
```



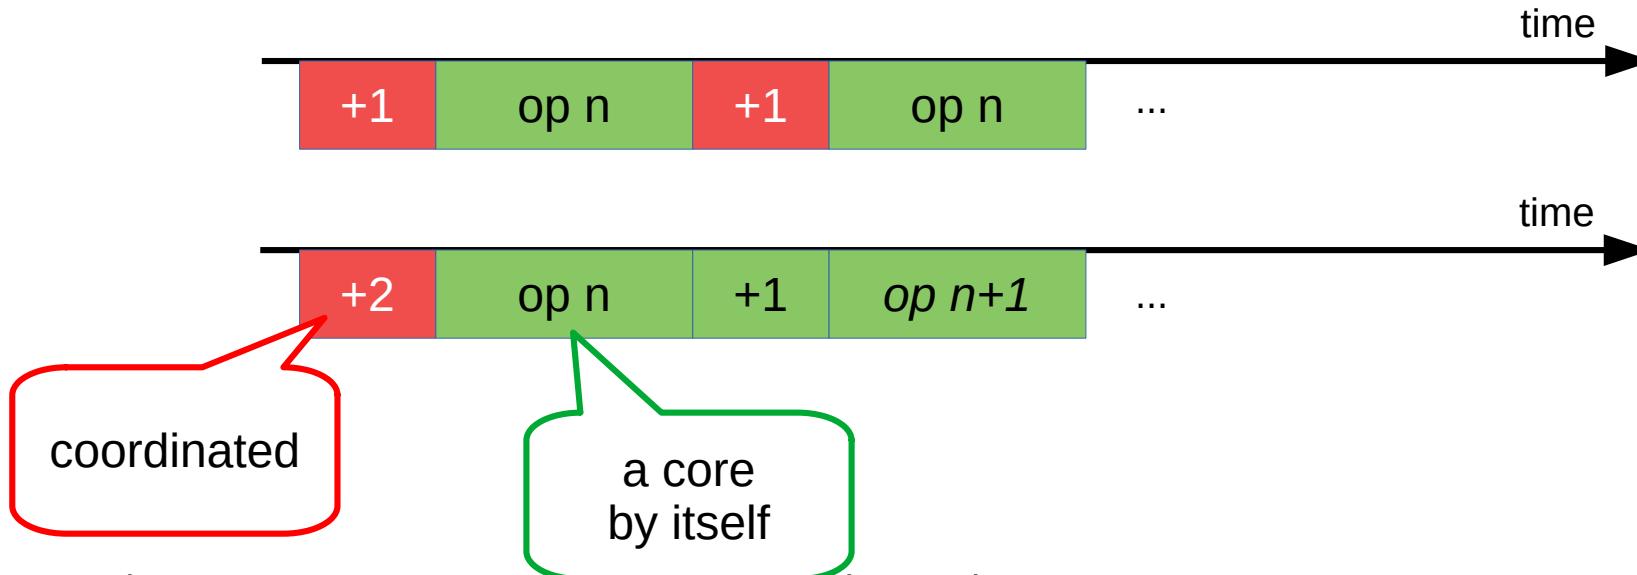
```
5      mov    eax, DWORD PTR [r
6      imul   eax, eax
7      mov    DWORD PTR [rbp-8]
8      mov    eax, DWORD PTR [rbp-12]
9      add    eax, eax
10     mov   DWORD PTR [rbp-12]
11     mov    DWORD PTR [rbp-4]
12     mov    eax, DWORD PTR [r
13     sub    eax, DWORD PTR [r
14     test   eax, eax
15     jns    .L2
16     mov    eax, DWORD PTR [r
17     add    DWORD PTR [rbp-4]
18     .L2:
19     mov    edx, DWORD PTR [r
20     mov    eax, DWORD PTR [r
21     add    edx, eax
22     mov    eax, DWORD PTR [r
```

Distributed

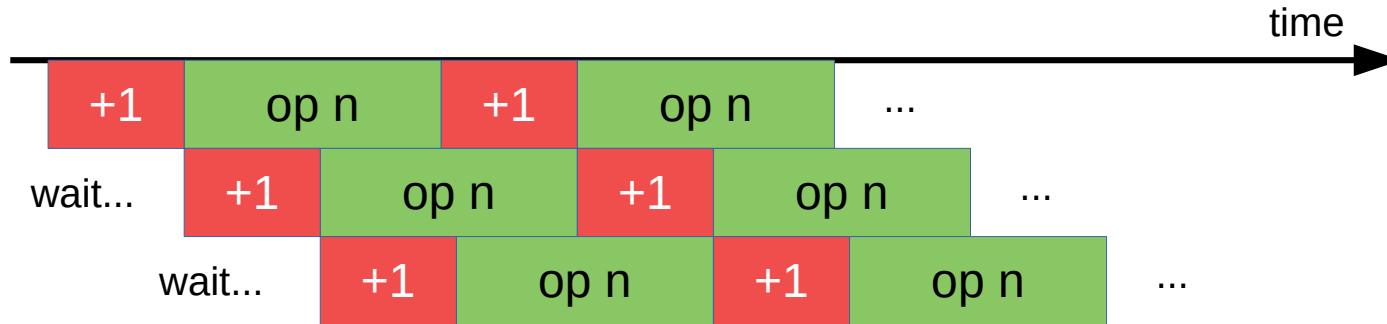


Coordination overhead

- Splitting a task incurs in coordination overhead
- Consider two versions of a chunked vector operation:
 - Get chunk of size 1, execute
 - Get chunk of size 2, execute one and the other

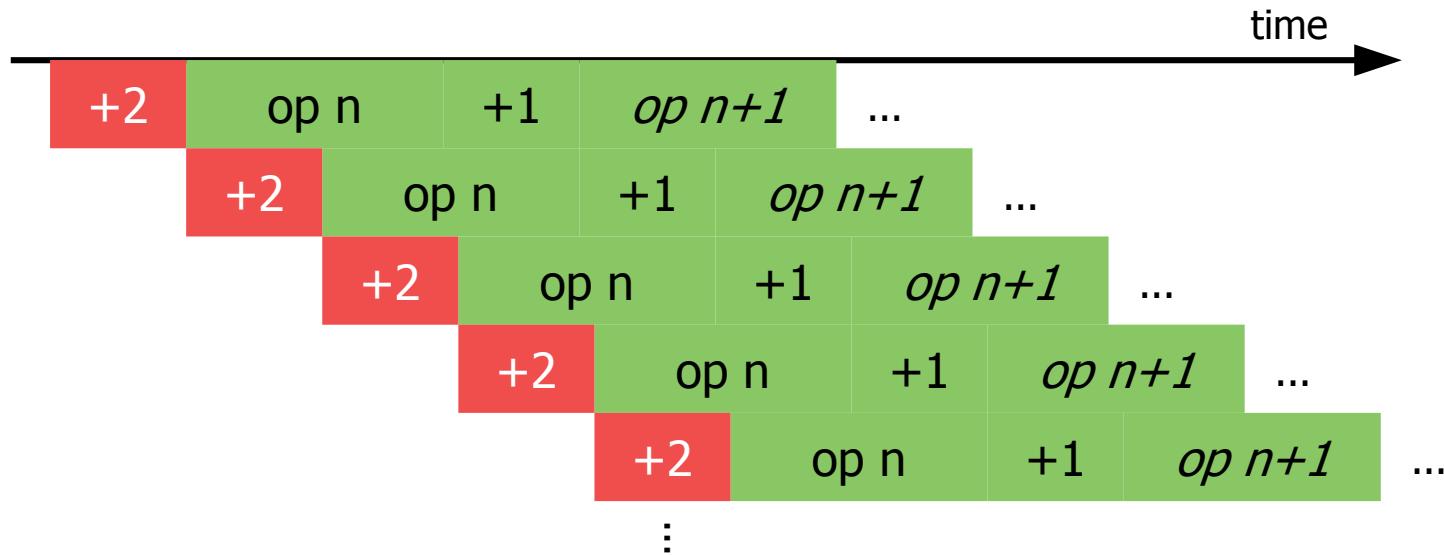


Coordination overhead



- Eventually, at least one core is blocked waiting for coordination

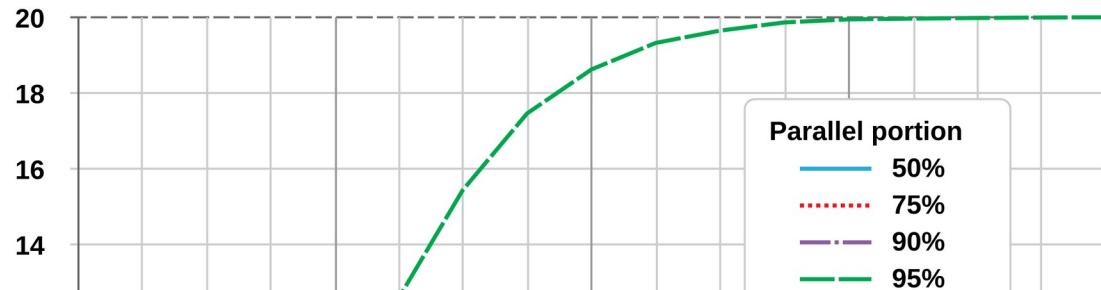
Coordination overhead



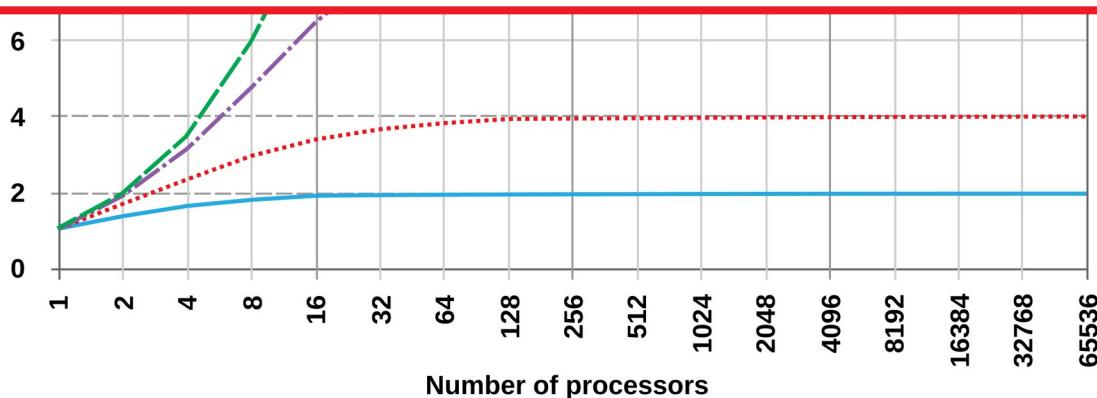
- Reducing the contention on coordination improves performance, even if doing the same work!

Amdahl's Law

Amdahl's Law



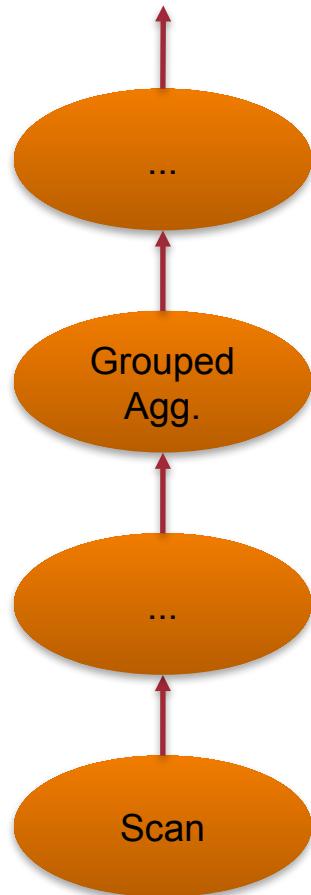
Key Issue:
How much time is used for coordination



Parallel execution

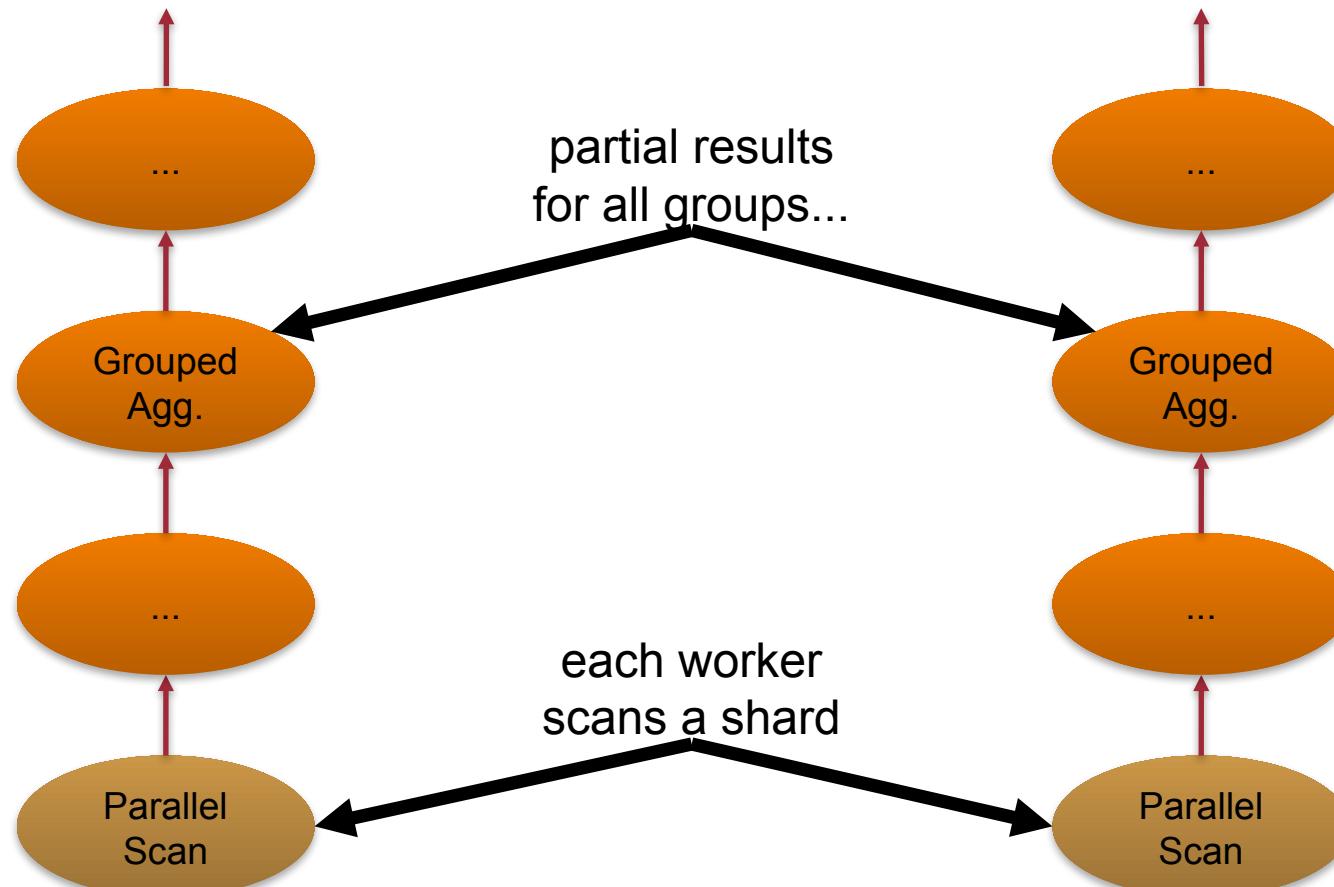
- Naive approach:
 - Split data in “shards”
 - Run an separate instance of the plan in each shard
 - Collect all results
- Sharding by:
 - Adding an additional filter to each scan when data is shared
 - Using the locally stored data when data is distributed

Example



```
select a, avg(b) from t  
where c > 0  
group by a  
having avg(b) > 0
```

Example: Naive parallel execution

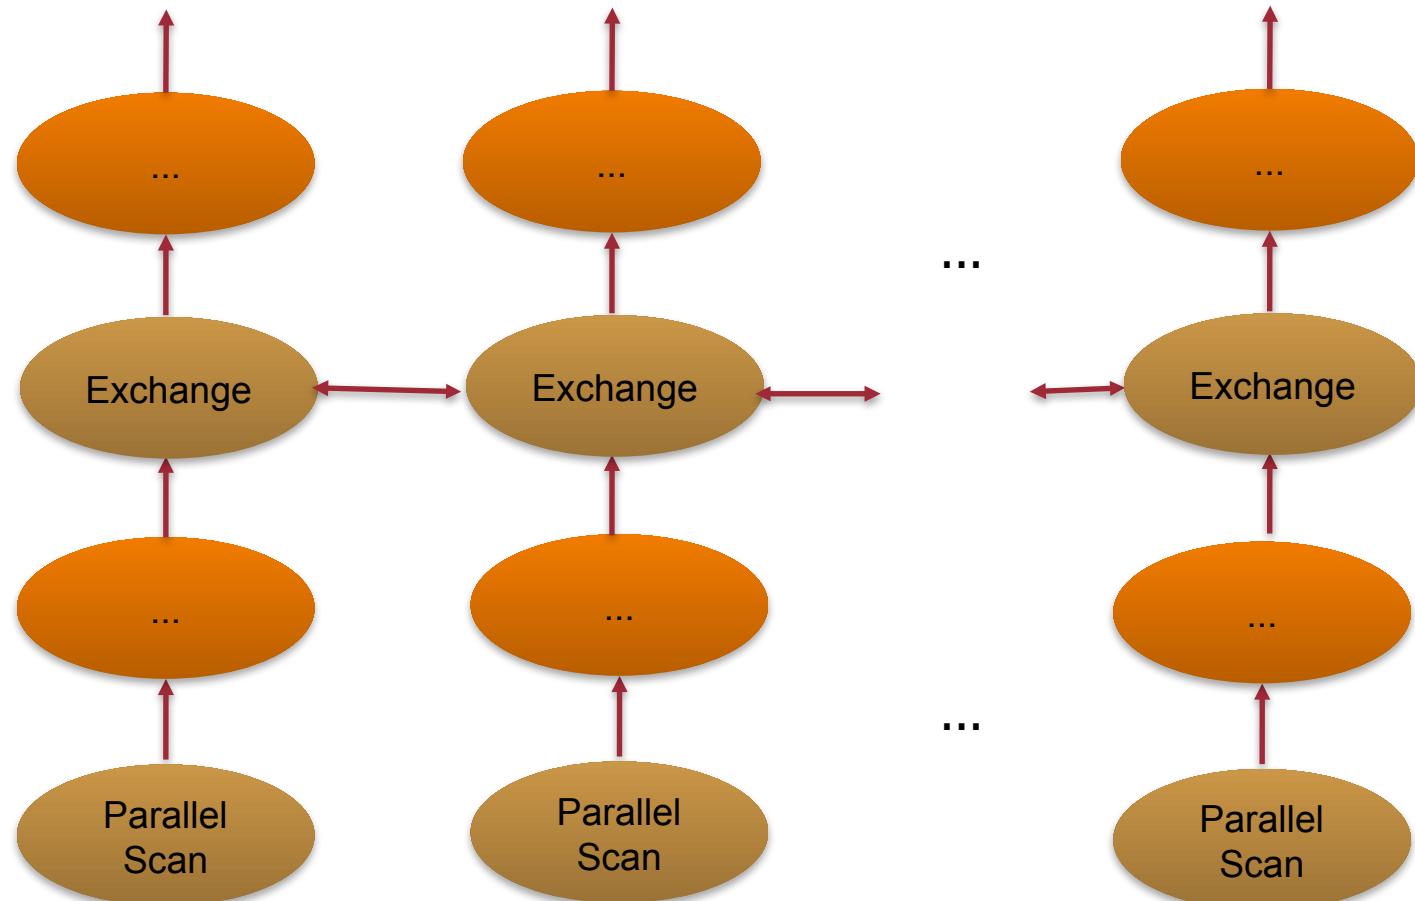


Volcano “Exchange” operator

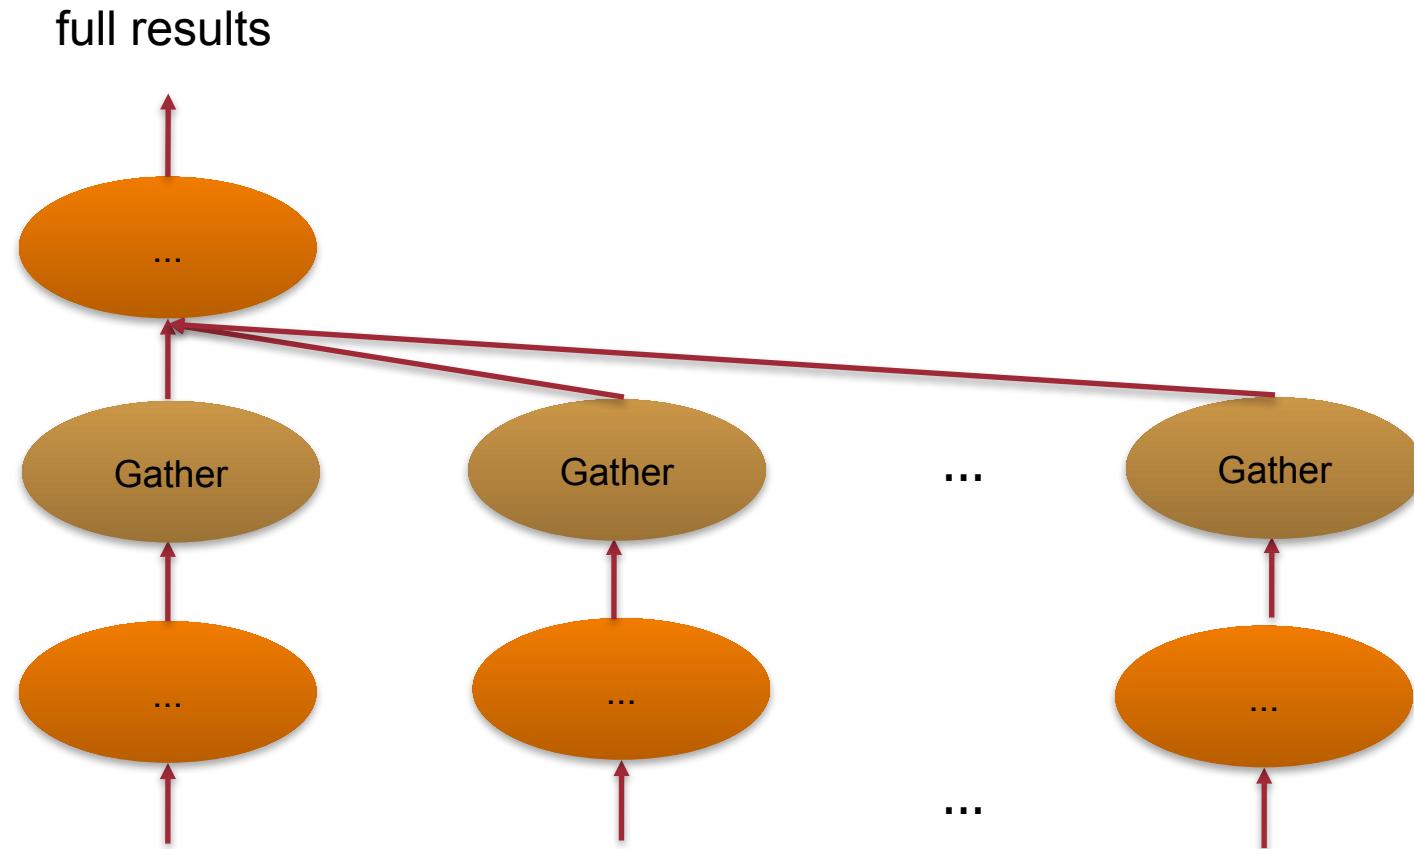
- Extends the iterator model to multi-core and distributed settings:
 - Contains an asynchronous, thread-safe (or distributed) buffer
 - Consumes data from multiple input iterators
 - Possibly each in a different copy of the plan
 - Can be consumed by multiple output iterators
 - Possibly each in a different copy of the plan
- Configurable rule to route tuples from inputs to outputs
- Different modes:
 - only one consumer → a.k.a. “gather”
 - respects order → a.k.a. “merge”
 - copies to all consumers → a.k.a. “broadcast”

Example: “Exchange” operator

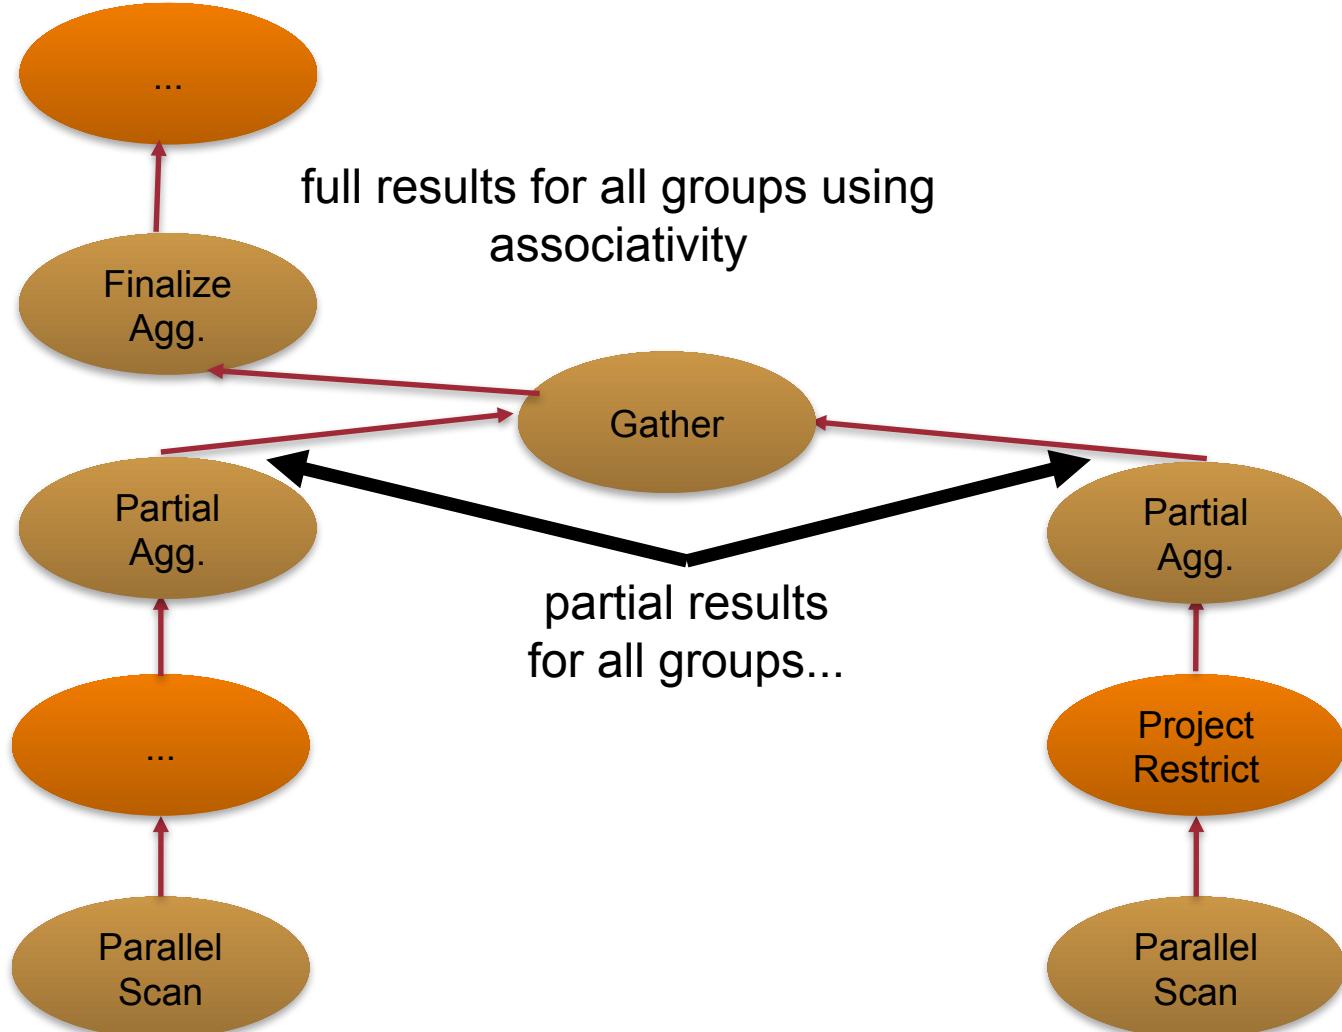
full results for some groups



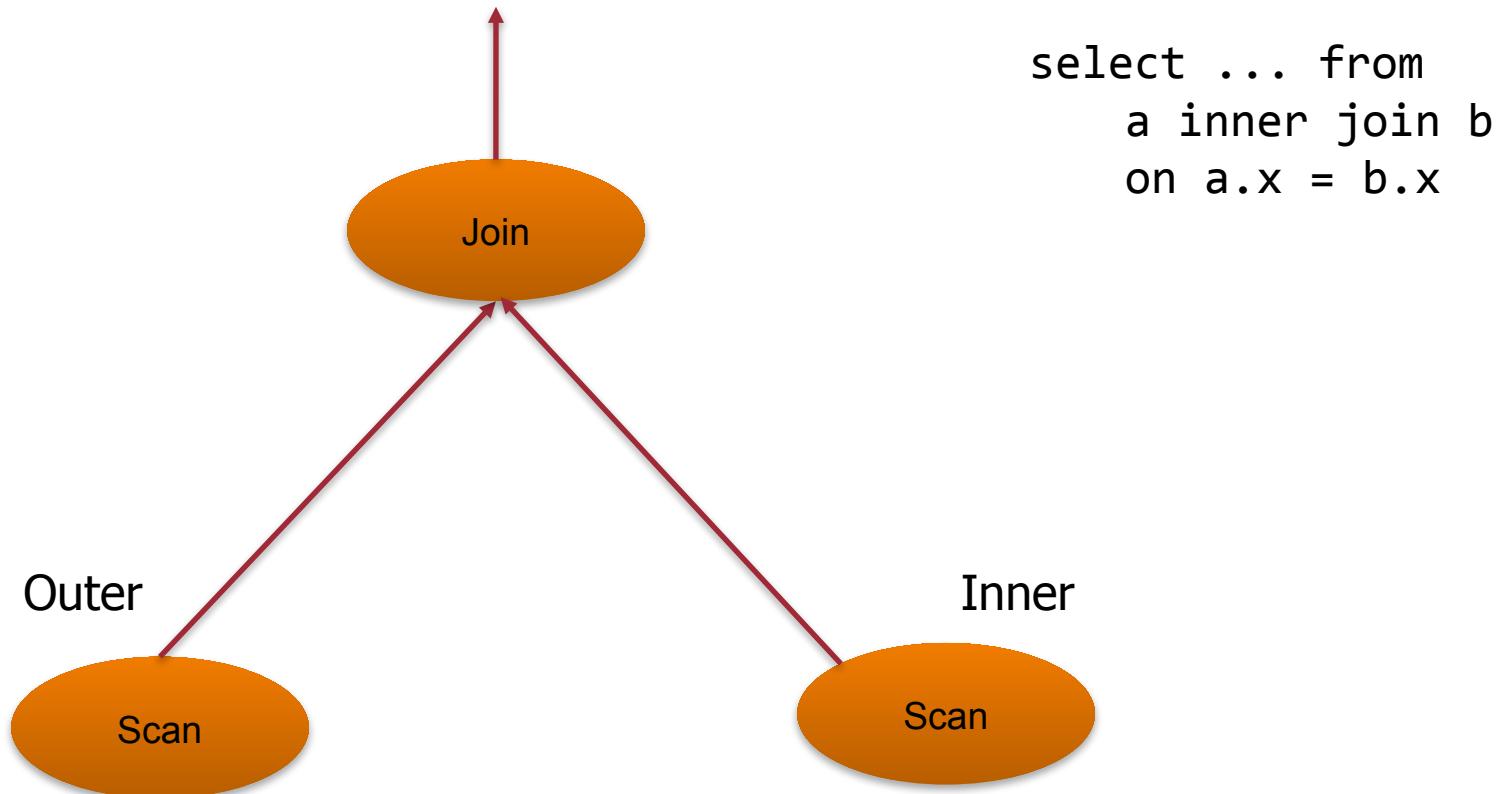
Example: “Exchange” operator



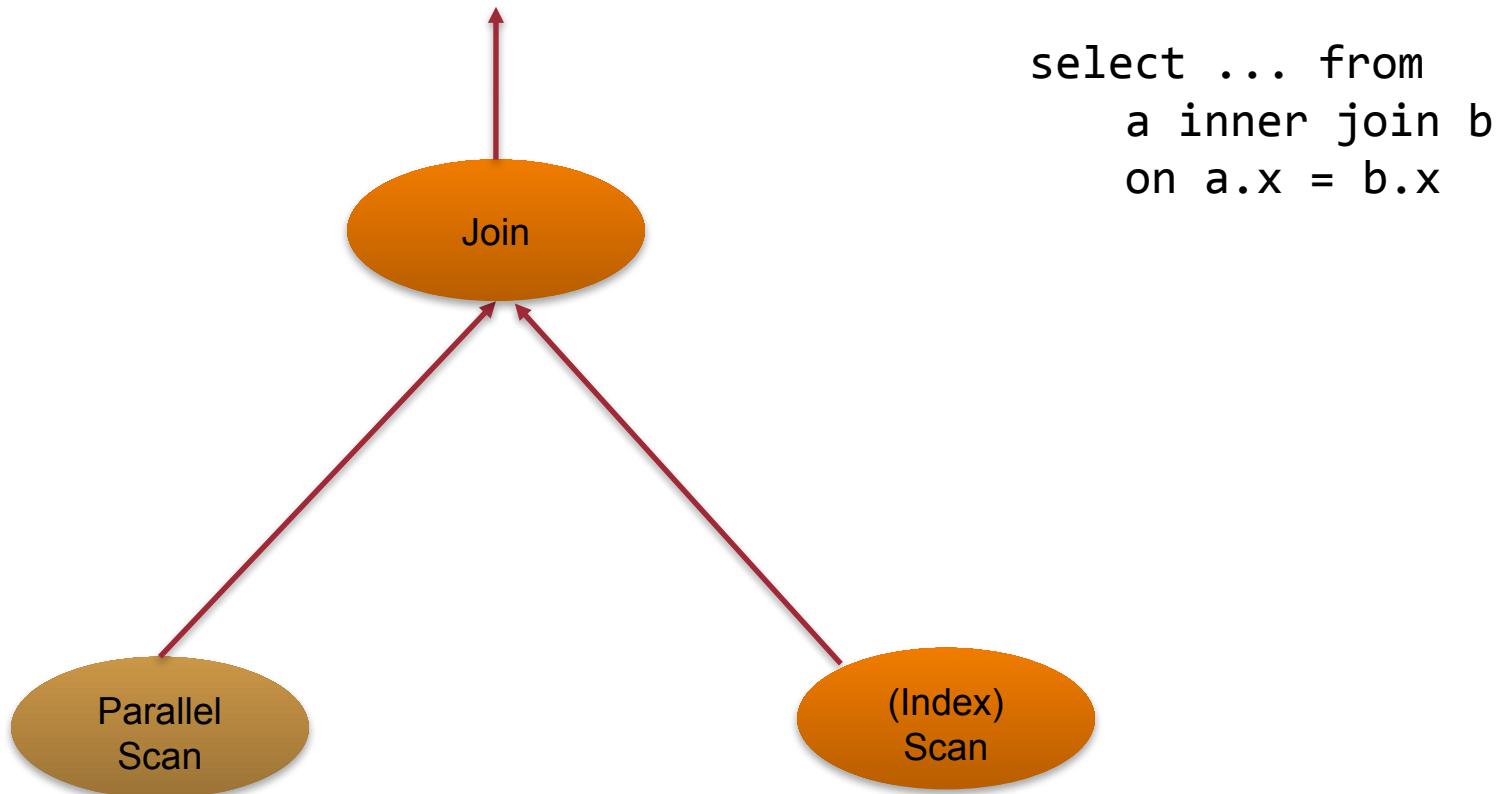
Optimized aggregation



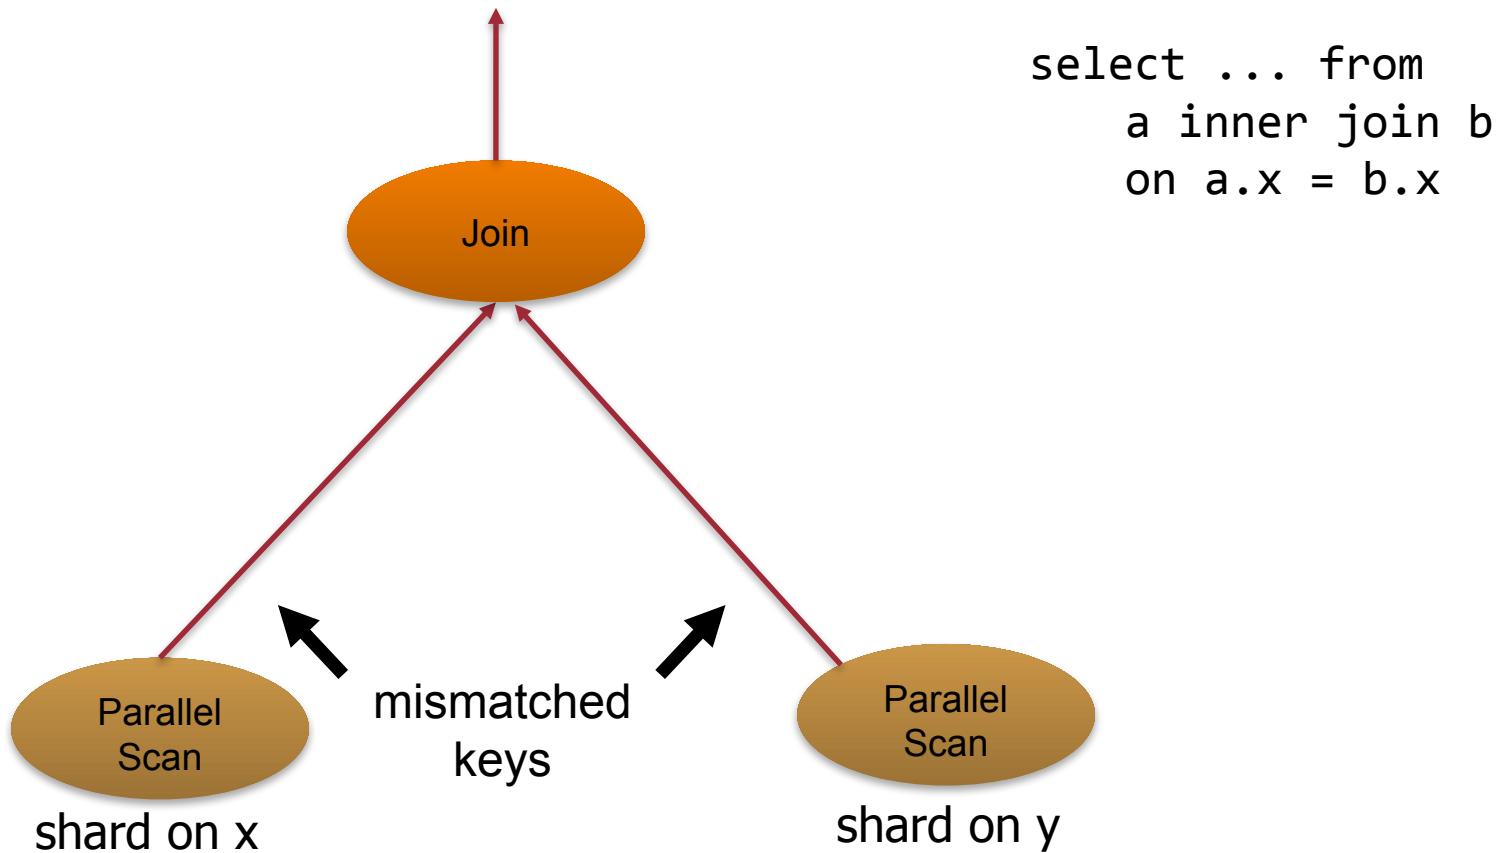
Parallel execution: Join



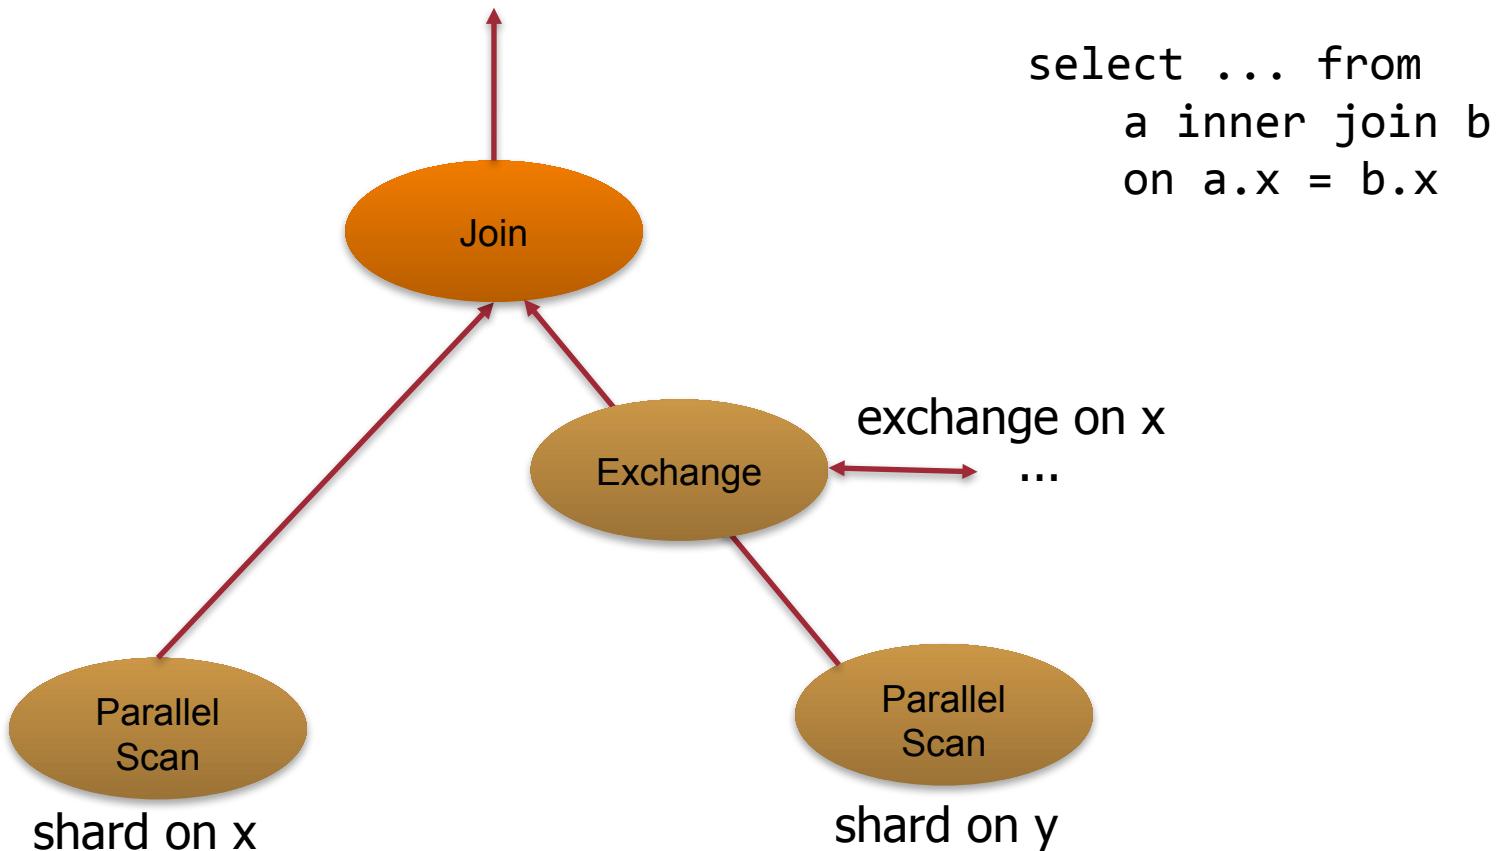
Parallel execution: NLJoin



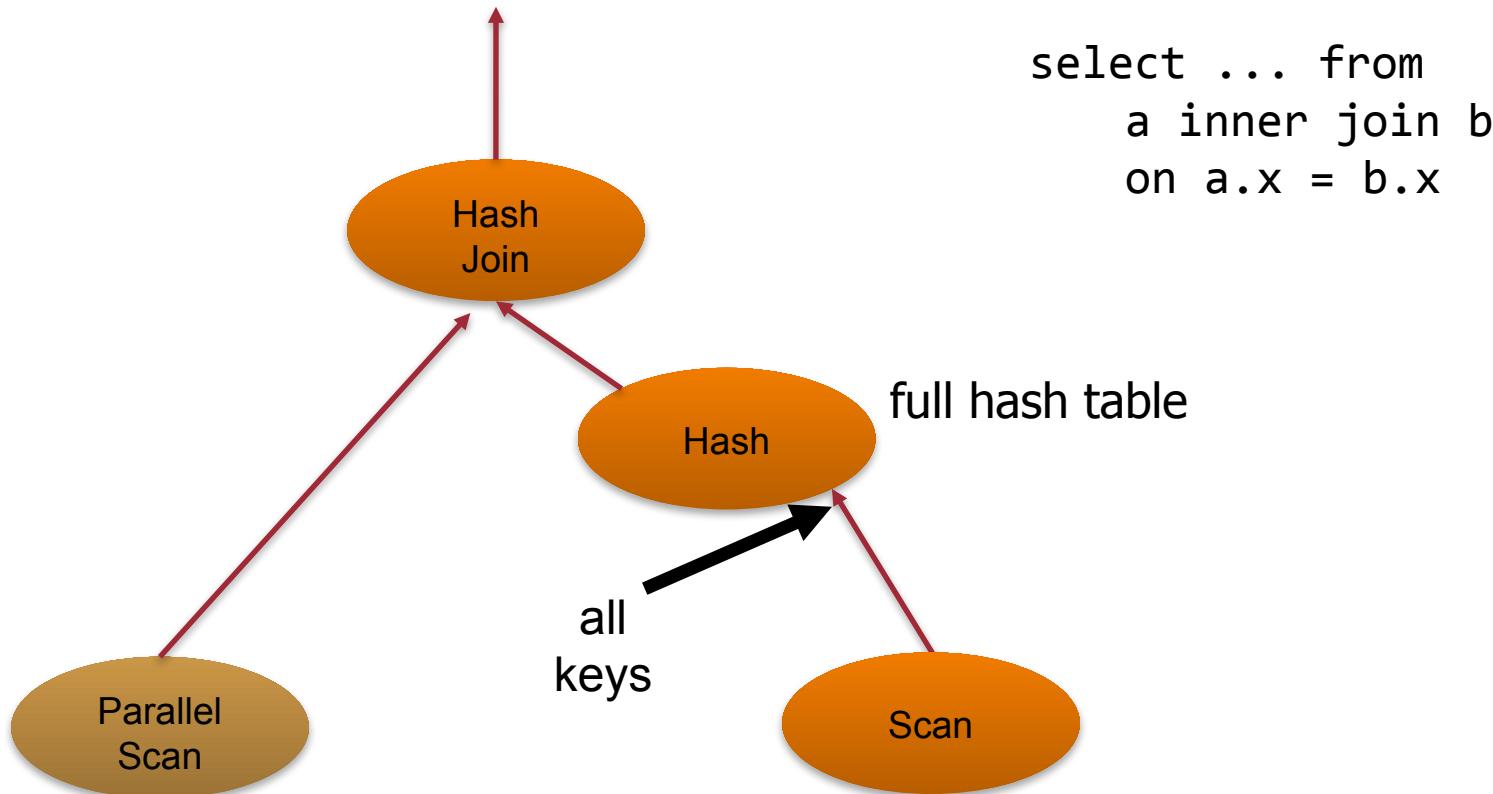
Parallel execution: MergeJoin



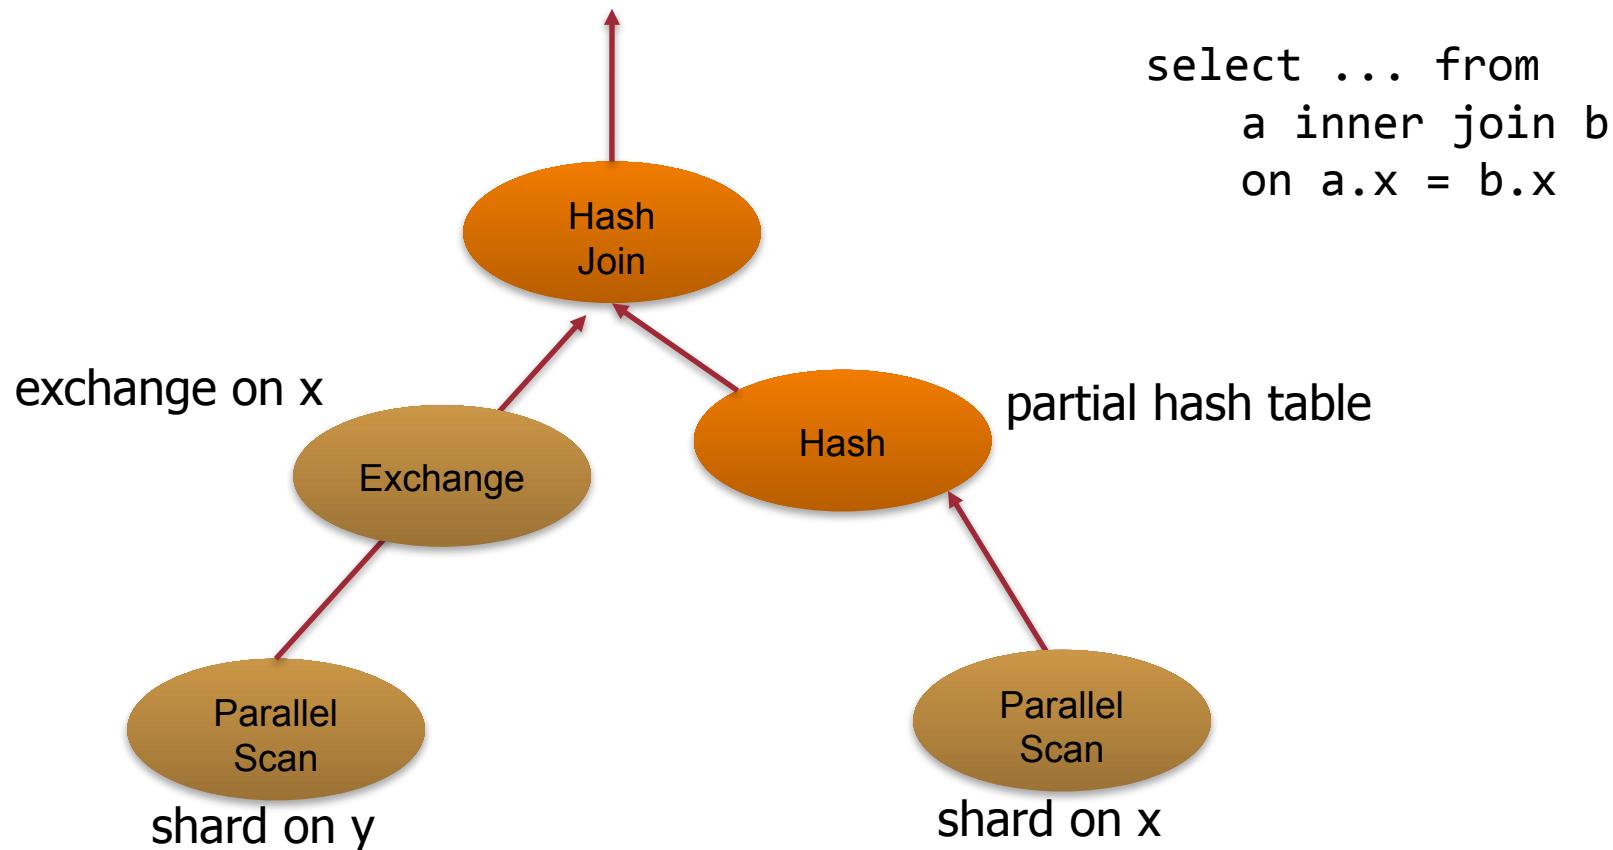
Parallel execution: Join



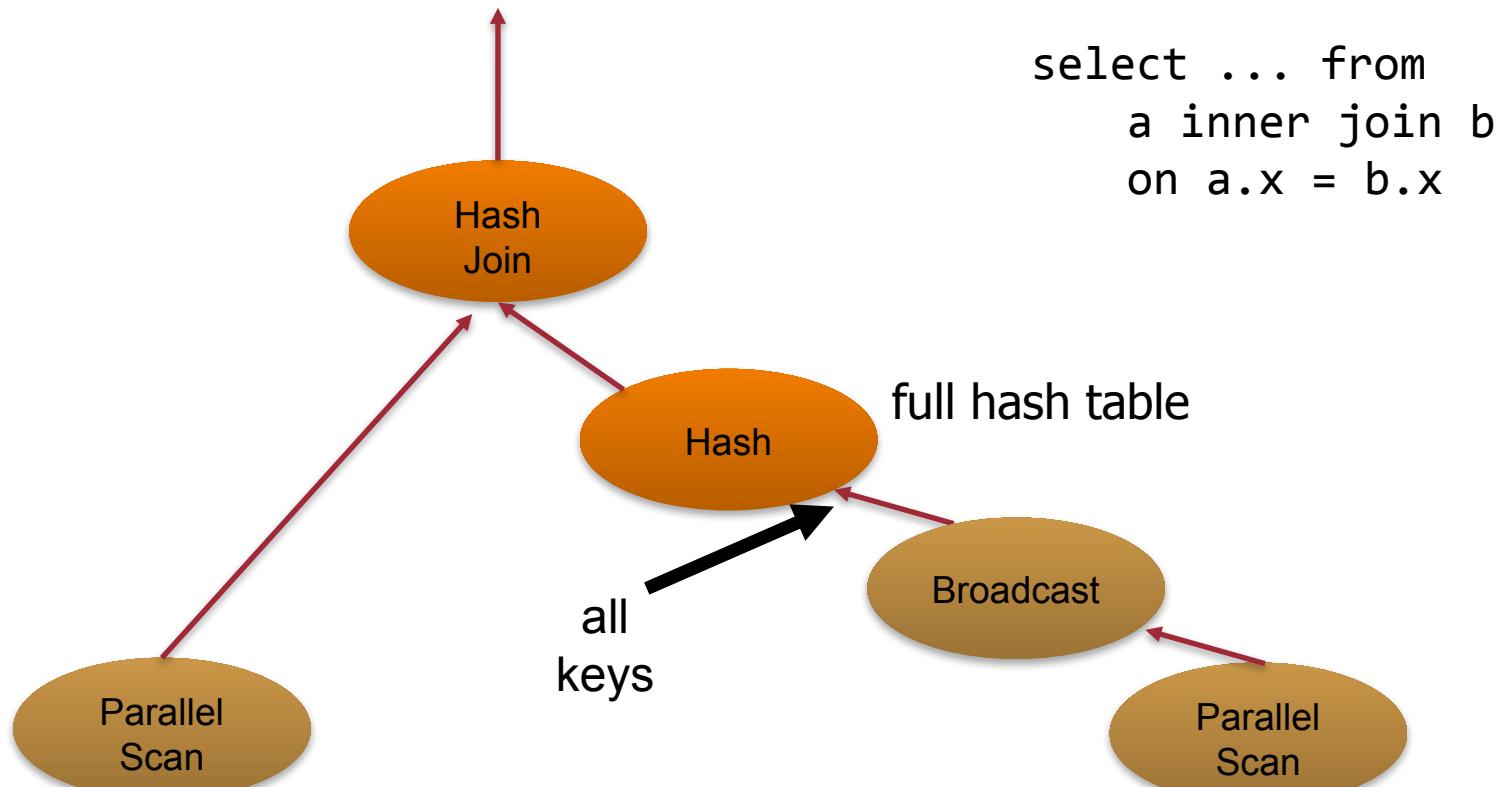
Parallel execution: Hash Join (shared)



Parallel execution: Hash Join (distributed)

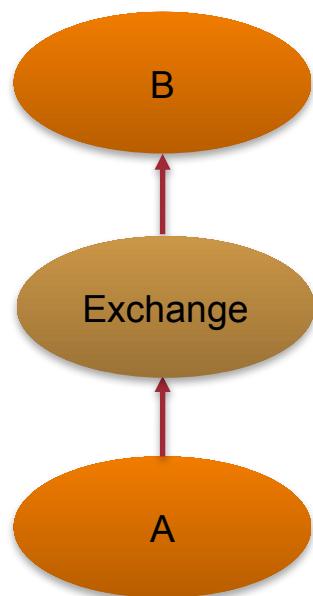


Parallel execution: Hash Join (distributed)



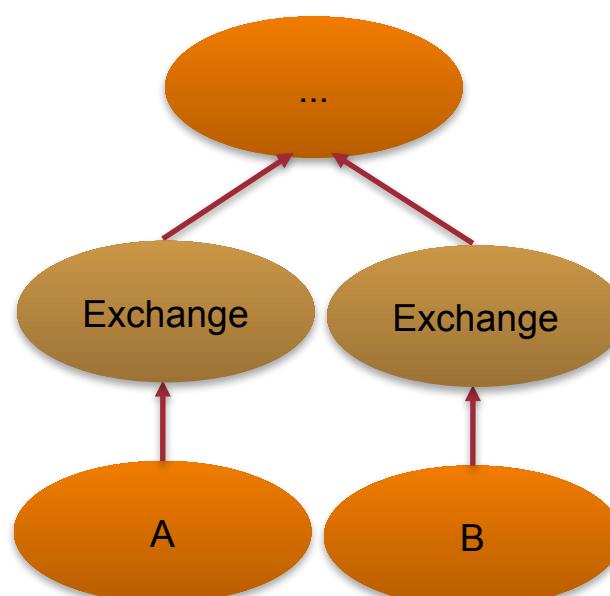
Consequences of Volcano “Exchange”

Vertical parallelism

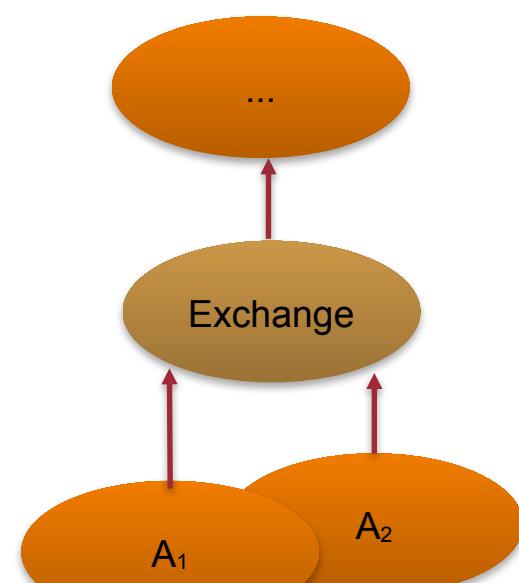


inter-operator

Horizontal parallelism



intra-operator



Best option!

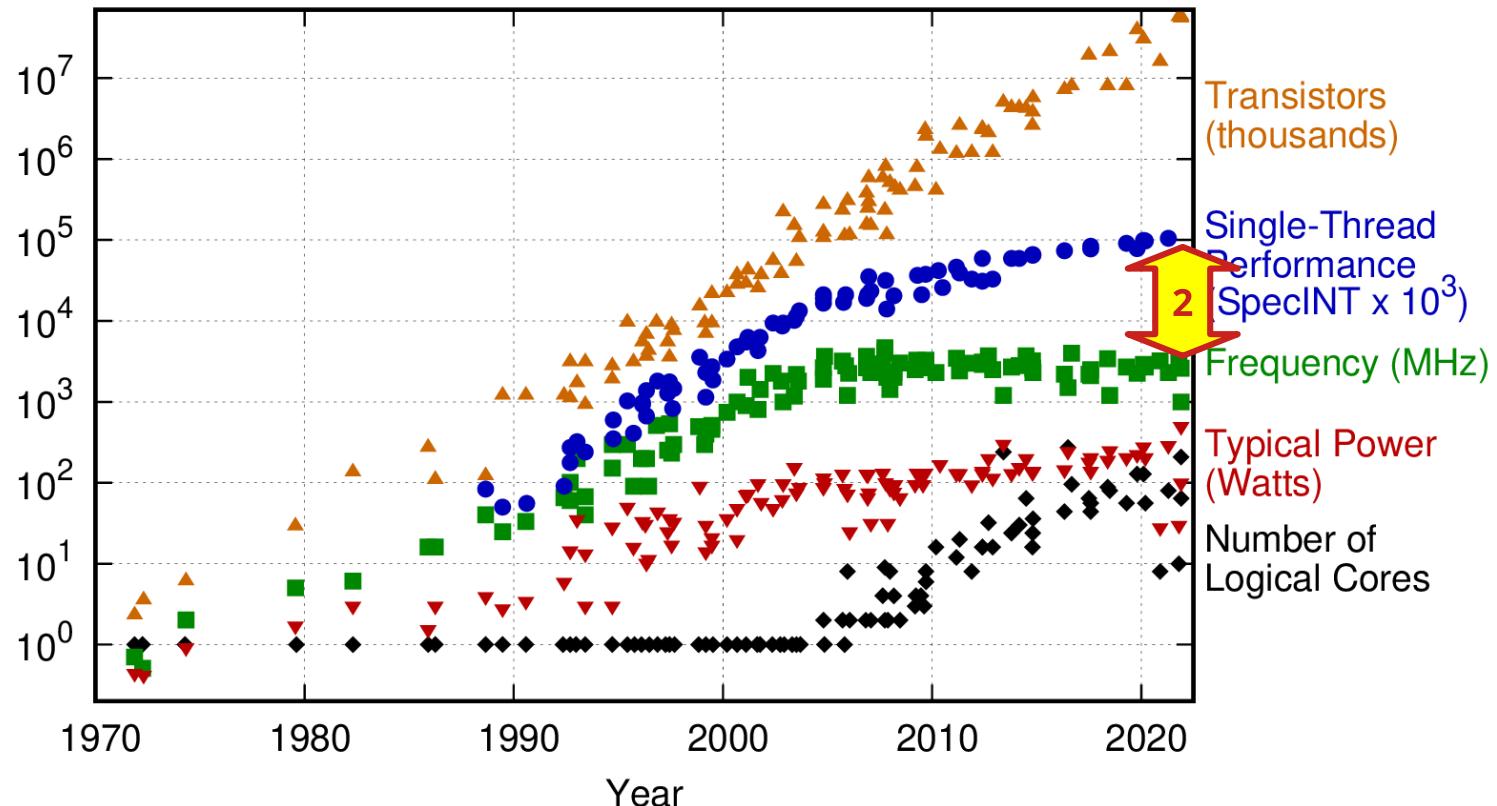
Consequences of Volcano “Exchange”

- Can easily retrofit existing database systems
- Parallelism is configured top-down, but optimal parallelism is discovered bottom-up
 - Data distribution
- Adequate for a small number of cores



Instruction parallelism

50 Years of Microprocessor Trend Data

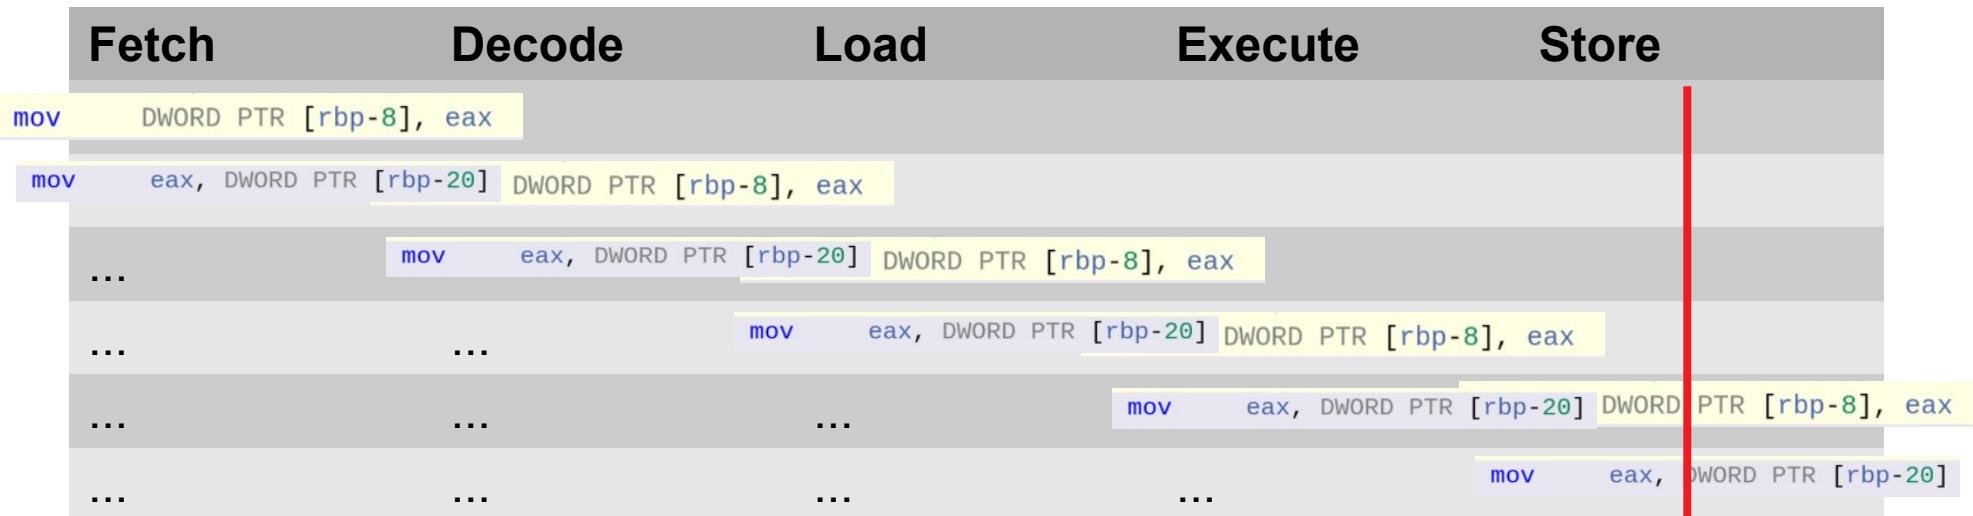


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp

Source <https://github.com/karlrupp/microprocessor-trend-data>

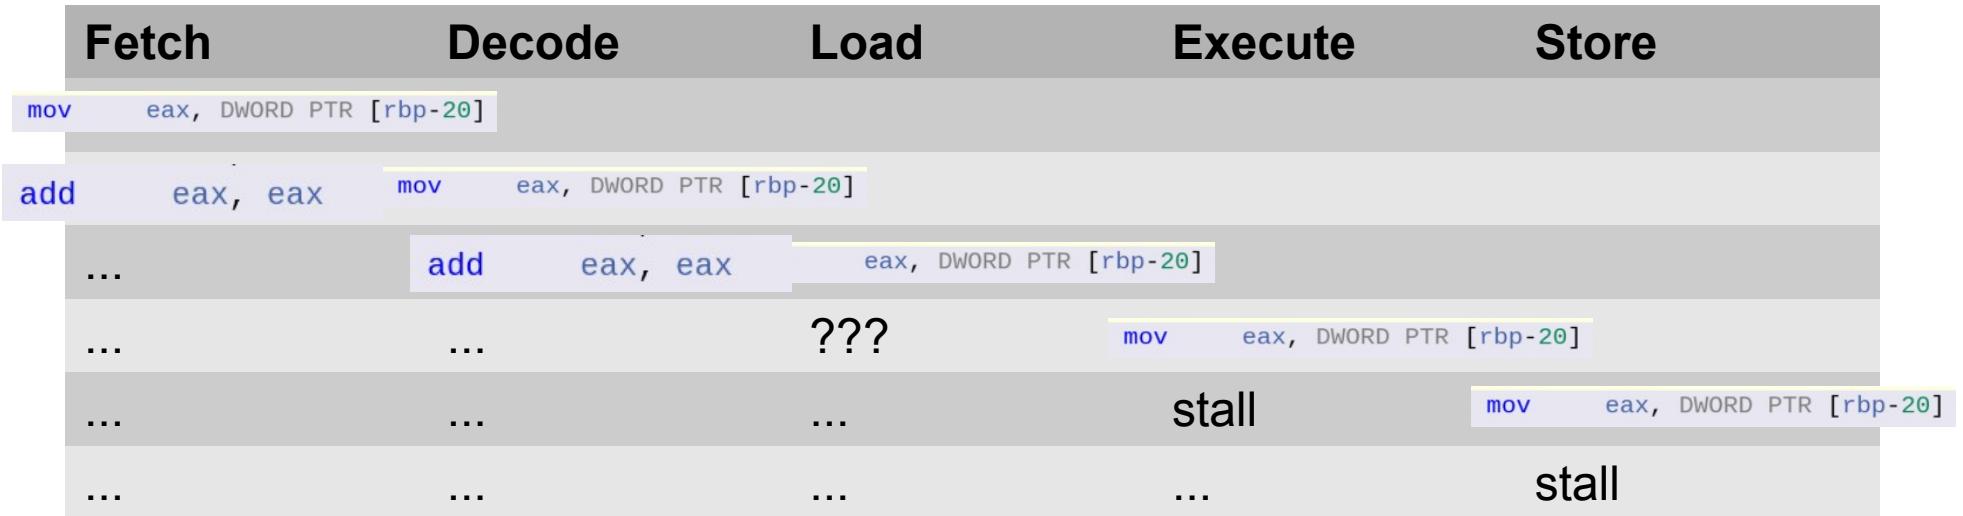
Pipelining

instruction latency = 5 cycles



throughput = 1 instruction / cycle

Pipelining



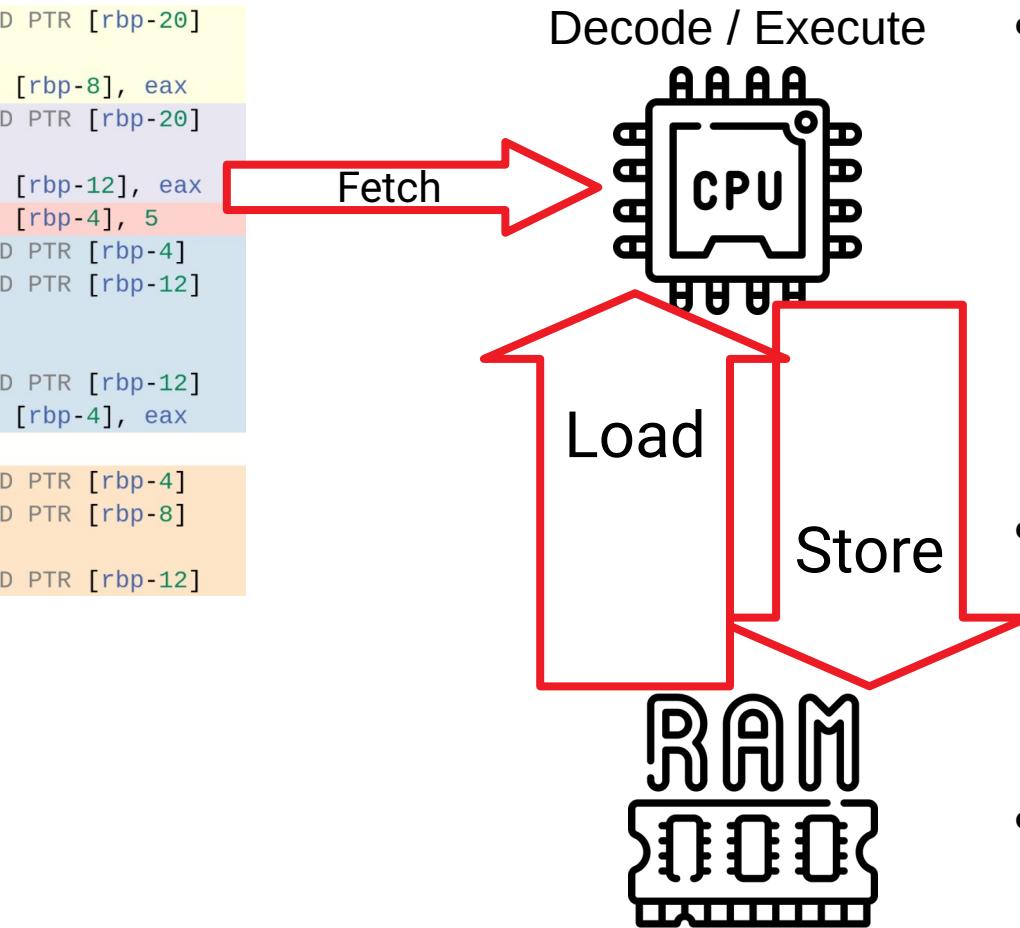
- Data dependency:
 - Trying to load a value that has not yet been computed

Pipelining

Fetch	Decode	Load	Execute	Store
jns .L2				
???	jns .L2			
stall	stall	jns .L2		
stall	stall	stall	jns .L2	
stall	stall	stall	stall	jns .L2
.L2:	mov edx, DWORD PTR [rbp-4]	stall	stall	stall

- Control flow dependency:
 - Cannot predict the next instruction

SIMD



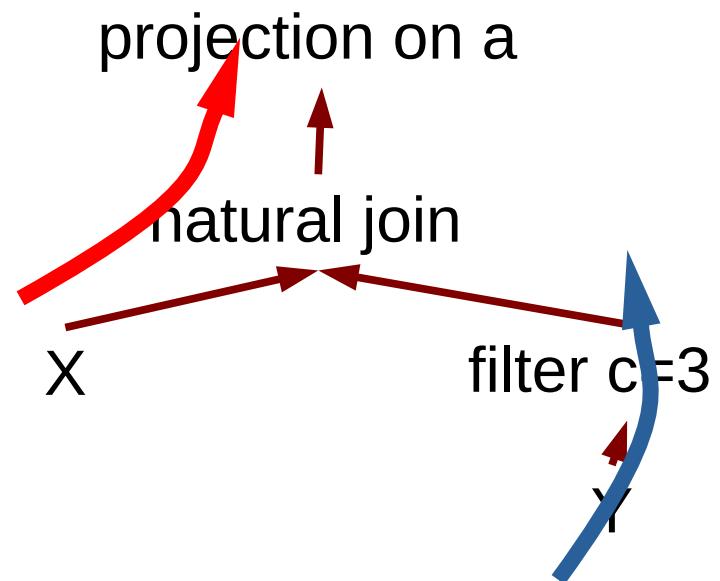
- Use wide registers that can fit vectors instead of scalars:
 - Example: Intel AVX512 → 512 bits
 - 64 byte vector
 - 32 shorts
 - 16 ints
 - ...
- Load, execute, and store full vectors, or slices of vectors, in a single instruction
- Key technique in GPUs

Consequences of iteration

- Close to worst case scenario for data movement and parallelism! 😱
 - Poor locality → Impacts caching / NUMA
 - Short code segments interleaved with dereferencing through virtual pointers → Processor pipeline stalls
 - Computation on one value at a time → No SIMD
- Severely impacts analytical workloads!

Operator fusion

- Split the execution plan into pipelines
 - 1: scan Y → filter → hash
 - 2: scan X → join → projection
- Allows optimization step with code generation:
 - Optimized for pipelining and locality



Push model

- Data producer calls consumer:
 - for row in table:
 - transform(row)
- Large overhead if there is a buffer between every two operators
- Viable with operator fusion
 - Buffer at the end of each pipeline

Chunked data

- Iterate over “chunks”:
 - Records → Records of arrays
- Advantages of columnar layout: SIMD

b	c
aaa	1
bbb	2
bbb	3
ccc	3
ddd	4

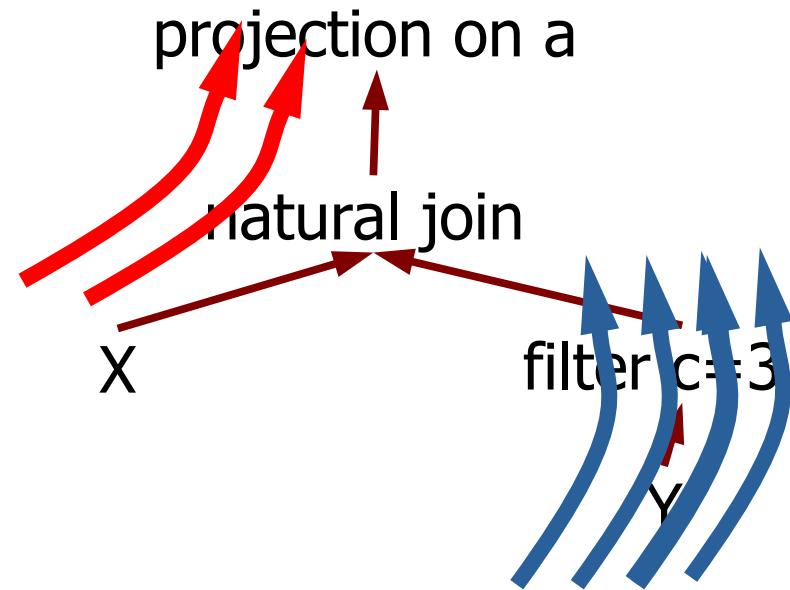
b	c
aaa	1
bbb	2
bbb	3
ccc	3
ddd	4

Chunked with operator fusion

- For each pipeline, generate a fused operator:
 - Outer loop waits for a chunk
 - Inner loop computes transformed chunk
 - Asynchronously delivers chunk
- Compile and optimize each pipeline:
 - Generate native code
 - Unroll inner loop, using SIMD instructions

Chunked with operator fusion

- Dynamically schedule multiple instances of each pipeline:
 - Depending on availability of input
 - Demand for output
 - Available resources



Consequences

- Ideal scenario for parallel processing:
 - Compiled fused operators → Simple and predictable control flow
 - Chunks → Good locality and SIMD
 - Pipeline scheduling → Dynamic multi-threading
 - Push model → Easier multi-thread scheduling



DuckDB

Conclusion

- There are a number of options for executing each query
- More options if we consider other data structures
- Varying performance:
 - Memory requirements
 - Number of iterations
 - Disk accesses
- What is the best one?
- How can it be discovered?