

Database Administration

José Orlando Pereira

Departamento de Informática
Universidade do Minho




Transaction

- A unit of work composed by individual read and write database operations
- Can be committed or rolled back

ACID Transactions

- Atomicity
- Consistency
- Isolation
- Durability



- The developer only cares about:
 - the logic of each transaction
 - running alone
 - in a perfect world

Atomicity vs. durability

- How to ensure both upon failure?
 - A single transaction might update multiple disk sectors
 - Upon reboot, any subset of pending sectors might have been updated

Atomicity vs. durability

- Durability without atomicity is easy:
 - Update database files before acknowledging commit
- Atomicity without durability is also easy:
 - Upon reboot, erase the database!

Goal

- A sequence of atomic operations is not an atomic operation
- There is a general purpose technique for obtaining an atomic sequence of operations

Requirements

- Additional storage, written sequentially (a log)
- Storage operations:
 - Individual operations are atomic
 - It is known when individual operations are finished
- Later, we need to remove / reuse old log entries

Redo log

- First method: Re-Do log
- Based on “re-doing missing writes”

Redo log

- The application changes the memory cache:


Memory cache:



Backing store:



Log:



Redo log

- Dirty blocks are copied to the log followed by a commit marker:

Memory cache:



Backing store:



Log:



Redo log

- Eventually, dirty blocks are copied to the backing store:

Memory cache:



Backing store:



Log:




Redo log

- Recovery after restart:
 - Copy blocks of transactions with commit markers from the log to storage
 - Forward (tail to head)

Backing store:



Log:



Redo log truncation

- When can a log prefix be removed?
- Naive approach: Write all changes of committed transactions to disk
- How to keep track of which transactions modified each block?
- Cannot write a memory block to disk if it has been modified by a running transaction:
 - No-steal
- What if a block has been modified by two transactions: one finished, one running?

Log format

- Physical: The block itself
 - Clearly idempotent
 - Inefficient
- Logical: The command that changes the block
 - Harder to make it idempotent (INSERT?)
 - Space efficient
- “Physiological”:
 - Either one of the above, depending on the operation, the context, ...

Redo log summary

- The good:
 - The transaction can be committed without modifying the backing store (no force)
- The bad:
 - Modifications can only be written to backing store after the transaction begins committing (no steal)
 - Checkpoints are hard
 - Assumes that memory is large enough to hold all modifications

Undo log

- Second method: Un-Do log
- Based on “un-doing unwanted writes”

Undo log

- Original values of dirty blocks are copied to the log:

Memory cache:



Backing store:



Log:



Undo log

- Modified blocks are copied to the backing store:

Memory cache:



Backing store:



Log:



Dirty blocks can now
be evicted from cache!

Undo log

- A commit marker is inserted in the log:

Memory cache:



Backing store:



Log:



Undo log

- Recovery after restart:
 - Copy back entries of transactions without commit markers
 - Backward (head to tail)

Backing store:



Log:



Undo log truncation

- When can a log prefix be removed?
- Can always write all memory to disk, even while there are running transactions
 - Steal

Undo log truncation

- The log must still be kept due to a possible very old running transaction


Memory cache:



Backing store:



Log:




Undo log truncation

- Waiting for all transactions to finish means:
 - Disk is now consistent
 - Log is no longer required
- Can we account for running transactions and avoid stopping the system?

Undo log truncation

- Non-quiescent (a.k.a. fuzzy) checkpoint:
 - Write to log “start checkpoint” marker with a list of running transactions $\langle T_1, \dots, T_n \rangle$
 - Wait until $\langle T_1, \dots, T_n \rangle$ have finished
 - Implicitly writes all modified memory blocks
 - Other transactions may start in between
 - Write “end checkpoint” marker


Log:



Undo log truncation

- Recovery to a complete checkpoint:
 - Search log until “start checkpoint”
 - No unfinished transactions exist before that


Log:



Undo log truncation

- Recovery to an incomplete checkpoint:
 - Search log until “start checkpoint”
 - Get list of possibly unfinished transactions
 - Continue searching the log until all found (likely to be close...)

Log:



Undo log

- The good:
 - Modifications can be written to backing store before entering commit (steal)
 - Lock granularity can be smaller than a block
 - Much easier checkpoints
 - Assumes only that the log is large enough to hold all modifications
- The bad:
 - Modifications must be written to backing store before entering commit (force)


Undo-Redo log

- Third method: Un-Do and Re-Do log
- Based on using both previous methods simultaneously


Undo-redo log

- Blocks can be changed after having been copied to the log (undo):

Memory cache:



Backing store:



Log:



Undo-redo log

- Before commit, new values are also copied to the log (redo):

Memory cache:



Backing store:



Log:



Undo-redo log

- A commit marker is inserted in the log:

Memory cache:



Backing store:



Log:



Undo-redo log

- Do both recovery procedures:


- Redo, tail to head
- Undo, head to tail

Does order matter?


Backing store:



Log:



Undo-redo log




- Order matters only if:
 - Unfinished t2 updates same item as finished t1
 - t2 reads x before t1 writes x
- Means that:
 - t1 and t2 would be concurrent transactions
 - Both would have a lock on x!

Undo-redo log summary

- If all log entries contain both undo and redo data:
 - A block can be always be copied to backing store (steal)
 - A transaction can always commit immediately (no force)

Backup


- A full backup may take hours
- How to copy the database while processing update transactions?



Backup

- Perform a checkpoint:
 - Marks start of backup
 - Make previous logs obsolete (if any)
- Copy database
 - Concurrently with normal operation
- Perform a second checkpoint:
 - Marks end of backup
- Copy logs between:
 - Start and end of backup markers

Case study: PostgreSQL



PostgreSQL

- Log method:
 - Redo log (a.k.a. WAL)
 - Undo log implicit in previous versions (updates never overwrite!)
- Log format:
 - Logical for most operations
 - Physical, the first time each block is modified after a checkpoint
- Background writer and periodical checkpoints
- On-line archiving and backup (a.k.a. PITR)

Conclusions

- Atomicity and durability ensured by undo-redo logging
 - Recovery
 - Checkpointing
- Backup adds one more level and different trade-offs, but the problem is fundamentally the same